

P4 Applied to a vSwitch Data Plane

Dan Daly

Intel Corp

June 4, 2015

How is a vSwitch Defined?

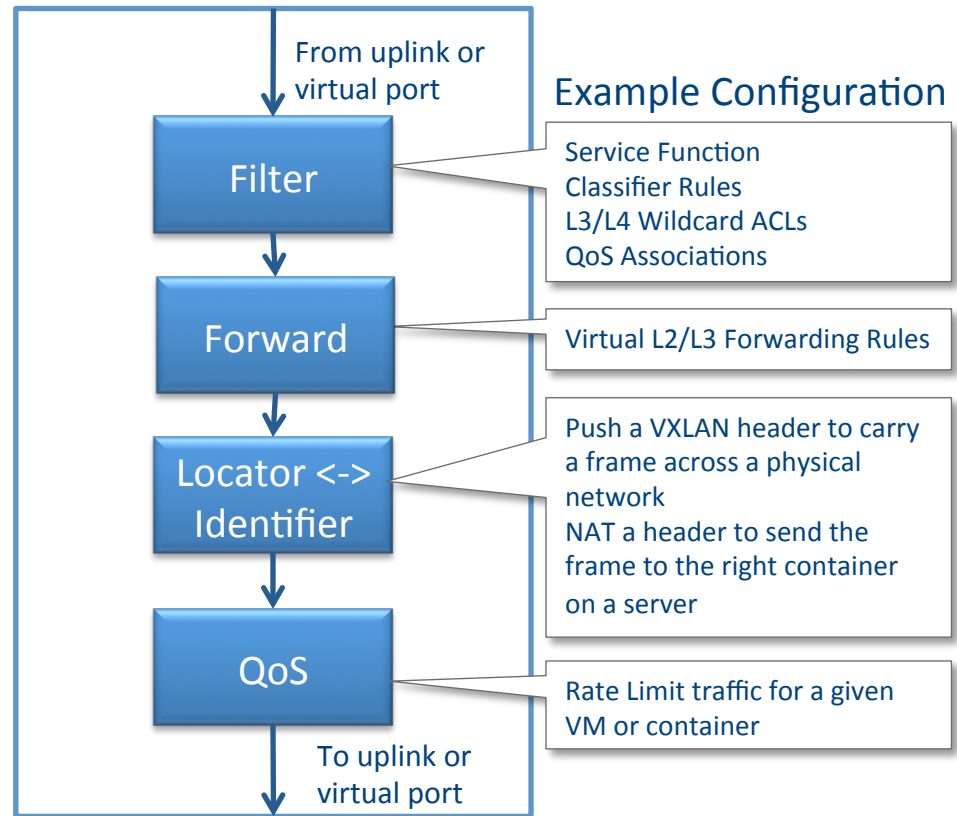
- literal: *Virtual Switch*, like a nexus 1000v
- de facto: Open vSwitch
- marketing: uses OpenFlow, enables SDN/NFV
- Kitchen sink: switches, routes, filters, tracks connections, ACLs, VPN, load balancer, snort, etc.

Why the question?

- Define the functionality of the vSwitch
- Optimize the vSwitch Data Plane on the server
- Problem Statement:
 - Virtualize the entire network
 - Virtual L2 and L3 domains
 - Virtual network services
 - High throughput, low latency & jitter

Abstract Virtual Switch Data Plane

- Requirements:
 - Virtual Forwarding
 - Locator/Identifier Separation
 - Filtering
 - QoS



Virtual Forwarding

- Locator: Marked on the packet to say where it should go
- Domain: Marked on the packet or in metadata as additional context
- Destination: a port, into a tunnel, a NAT, or into another pipeline (trap case)

{Locator, Domain} -> Destination

Locator/Identifier Separation

- Locator (Virtual) separate from Identifier (Physical)
- LISP, VXLAN, NVGRE, do this
- Address Translation can also do this

{Locator, Domain} -> Add(Identifier)

{Locator, Domain} -> Remove(Identifier)

- Irrespective of approach, the vSwitch operates on Locators, physical network can rely on Identifiers

Filtering

```
table filtering {  
  reads {  
    standard_metadata.ingress_port :  
      wildcard;  
    ethernet.dst_mac : wildcard;  
    ethernet.src_mac : wildcard;  
    ...  
    ip.src_ip : wildcard;  
    ip.dst_ip : wildcard;  
    l4.src_port : wildcard;  
    l4.dst_port : wildcard;  
    tcp.flags : wildcard;  
    vxlan.vni : wildcard;  
    nsh.nsp : wildcard;  
    nsh.nsi : wildcard;  
  }  
}  
  
actions {  
  count;  
  drop;  
  set_egress_port;  
  set_tunnel_id;  
  decap;  
  route_via_ecmp(group);  
  set.vlan1;  
  push_vlan(tpid, tci);  
  pop_vlan;  
  mirror,  
  sample(probability, actions);  
}
```

Enable simple state tracking

Match on metadata tags for service chaining / forward graphs

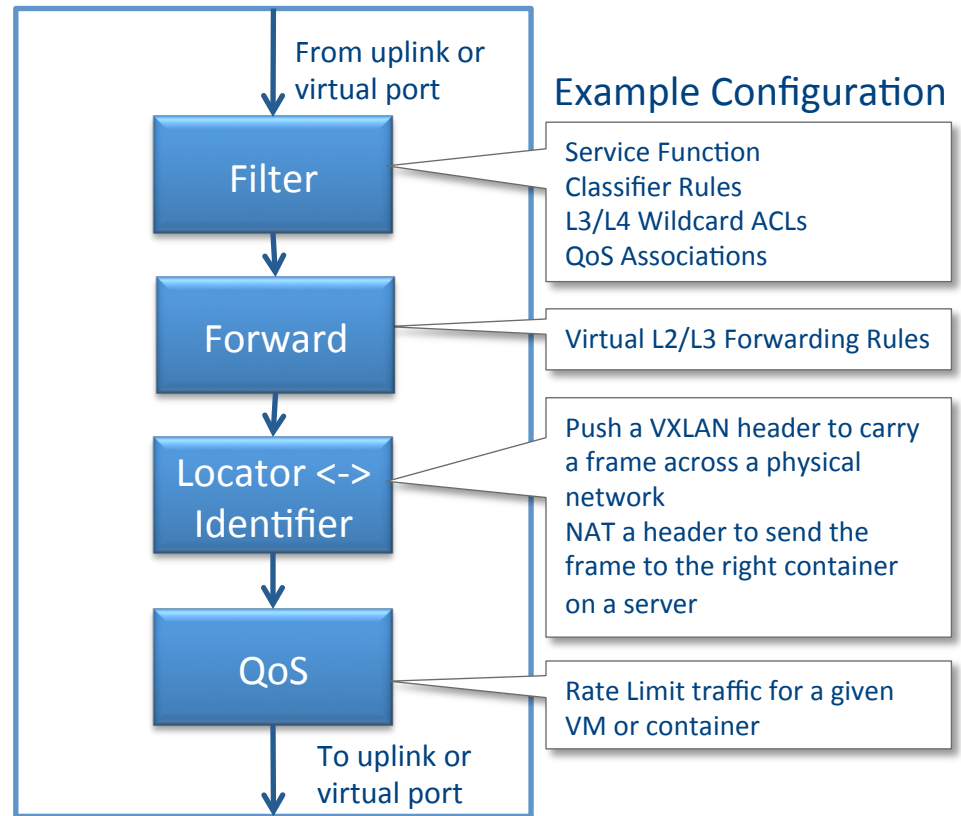
QoS

- Simplest Example: Each virtual port is given a virtual speed, so that each VM or container cannot exceed its virtual allotment

```
attribute {  
    type { configuration, per-port }  
    name { port.limit.bandwidth }  
    description { "Set a bandwidth limit for the  
given port. Note that the platform may restrict  
or even change the set of possible  
configurations" }  
    range { 1e6..100e9 }  
}
```

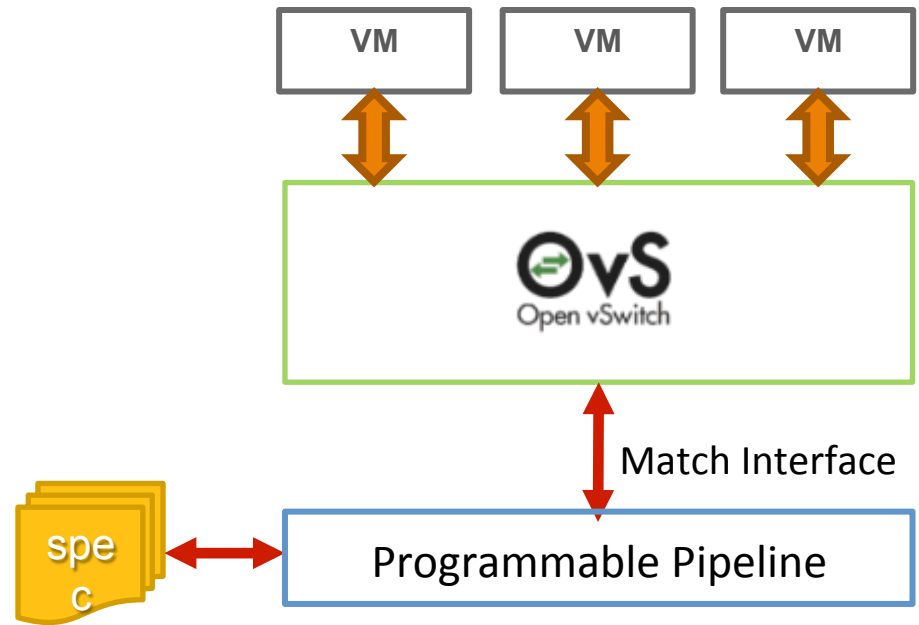

P4 Definition of a vSwitch Data Plane

- Concisely Defines:
 - Match Conditions
 - Actions Supported
 - Tables
 - Configuration
 - Headers
- Specifies application intent
 - “I want to filter on X”
 - “My tunnel looks like Y”



OvS Integration

- Adding a programmable data plane under OvS
- Similar to how the kernel data plane is exposed
- Converts OvS rules into the underlying data plane



<https://github.com/match-interface>

Summary

- P4 can be used to concisely define the requirements for a vSwitch Data Plane
 - What can the vSwitch do? Read the P4
- P4 could also be used to define future requirements
 - Example: Hooks for connection tracking
- Open source project to integrate a programmable pipeline underneath OvS:
<https://github.com/match-interface>