



XILINX

ALL PROGRAMMABLE™

P4 for an FPGA target

Gordon Brebner
Xilinx Labs
San José, USA

P4 Workshop, Stanford University, 4 June 2015


What this talk is about

- **FPGAs and packet processing languages**
- **Xilinx SDNet data plane builder featuring the PX language**
- **P4 and PX comparison**
- **P4-to-FPGA compilation using SDNet**
- **Protocol Independent Forwarding (PIF) project**
- **Summary**

FPGA: the “white box” hardware chip

➤ Attributes of Xilinx Ultrascale FPGAs in 2015:

Up to
2.5m
tiles



Tile is 6-input logic gate and 2 flip-flops
+ Embedded function blocks and memories
+ 100G Ethernet and 150G Interlaken interfaces

Local and long-distance wiring between components

All **programmable** by writing to memory

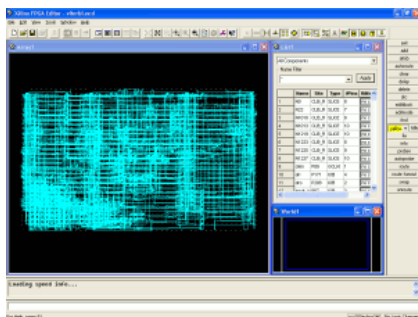
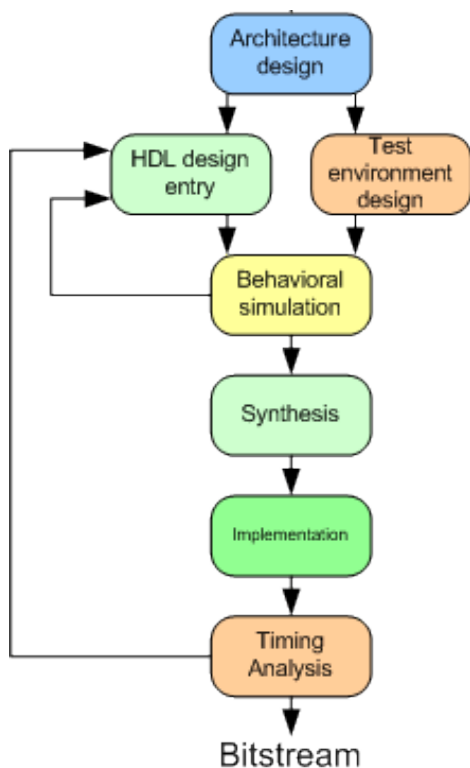
➤ Can now implement complex packet processing on single chip

- Open platform used by researchers: NetFPGA

➤ Beyond single chips:

- Multiple FPGAs, e.g., Corsica SDN data planes
- FPGA fast paths, CPU slow paths and control
- FPGA smart paths, ASIC dumb switches

Generic FPGA programming



- Chip design style experience
- Hardware Description Language (HDL)
- Cryptic results sometimes
 - Behavior and performance
- Enhancements (for ‘hardware guys’):
 - Libraries of blocks
 - Allow re-use and sharing
 - In HDL, or pre-synthesized
 - High-level synthesis
 - Typically superset of subset of C
 - Translated into HDL
- Abstraction needed for ‘software guys’
 - Current emphasis at Xilinx: **SDx environments**

Packet processing research in Xilinx Labs

➤ First generation (1 to 2 Gb/s line rates, 2001-2005)

- Initially, fine-grain Click as domain-specific language [DAC 2004]
- Then, coarse-grain Click plus HAEC state machine language [FPT 2004]

➤ Second generation (10 to 20 Gb/s line rates, 2005-2009)

- G language for packet processing functions ...
- ... plus Click (with extended semantics) for composition of functions
- Eventually published late in life [NetFPGA 2009, FPT 2009]

➤ Third generation (100+ Gb/s line rates, 2009-2013)

- PX language for packet processing functions and their composition
- First version published using PP as name [ANCS 2011]
- Enhanced version published using PX as name [IEEE Micro 2014]
- Productized as **Xilinx SDNet** in 2014

Xilinx SDNet Programmable Packet Processor (www.xilinx.com/sdnet)

➤ **Headline feature set, facilitated by FPGA white box target:**

100G+ packet rate
performance




Domain-specific
programming abstraction

```
class MPLS_TYPE {  
  struct{ label : 20,  
         cos   : 3,  
         sBit  : 1,  
         ttl   : 8 }  
  
  method next_header =  
    if (sBit == 0){  
      MPLS_TYPE;  
    } else {  
      ETH_TYPE;  
    }  
  
  method next_offset = size();  
  method earliest = 1;  
  method latest = 3;  
  method key_builder = (label)  
}
```

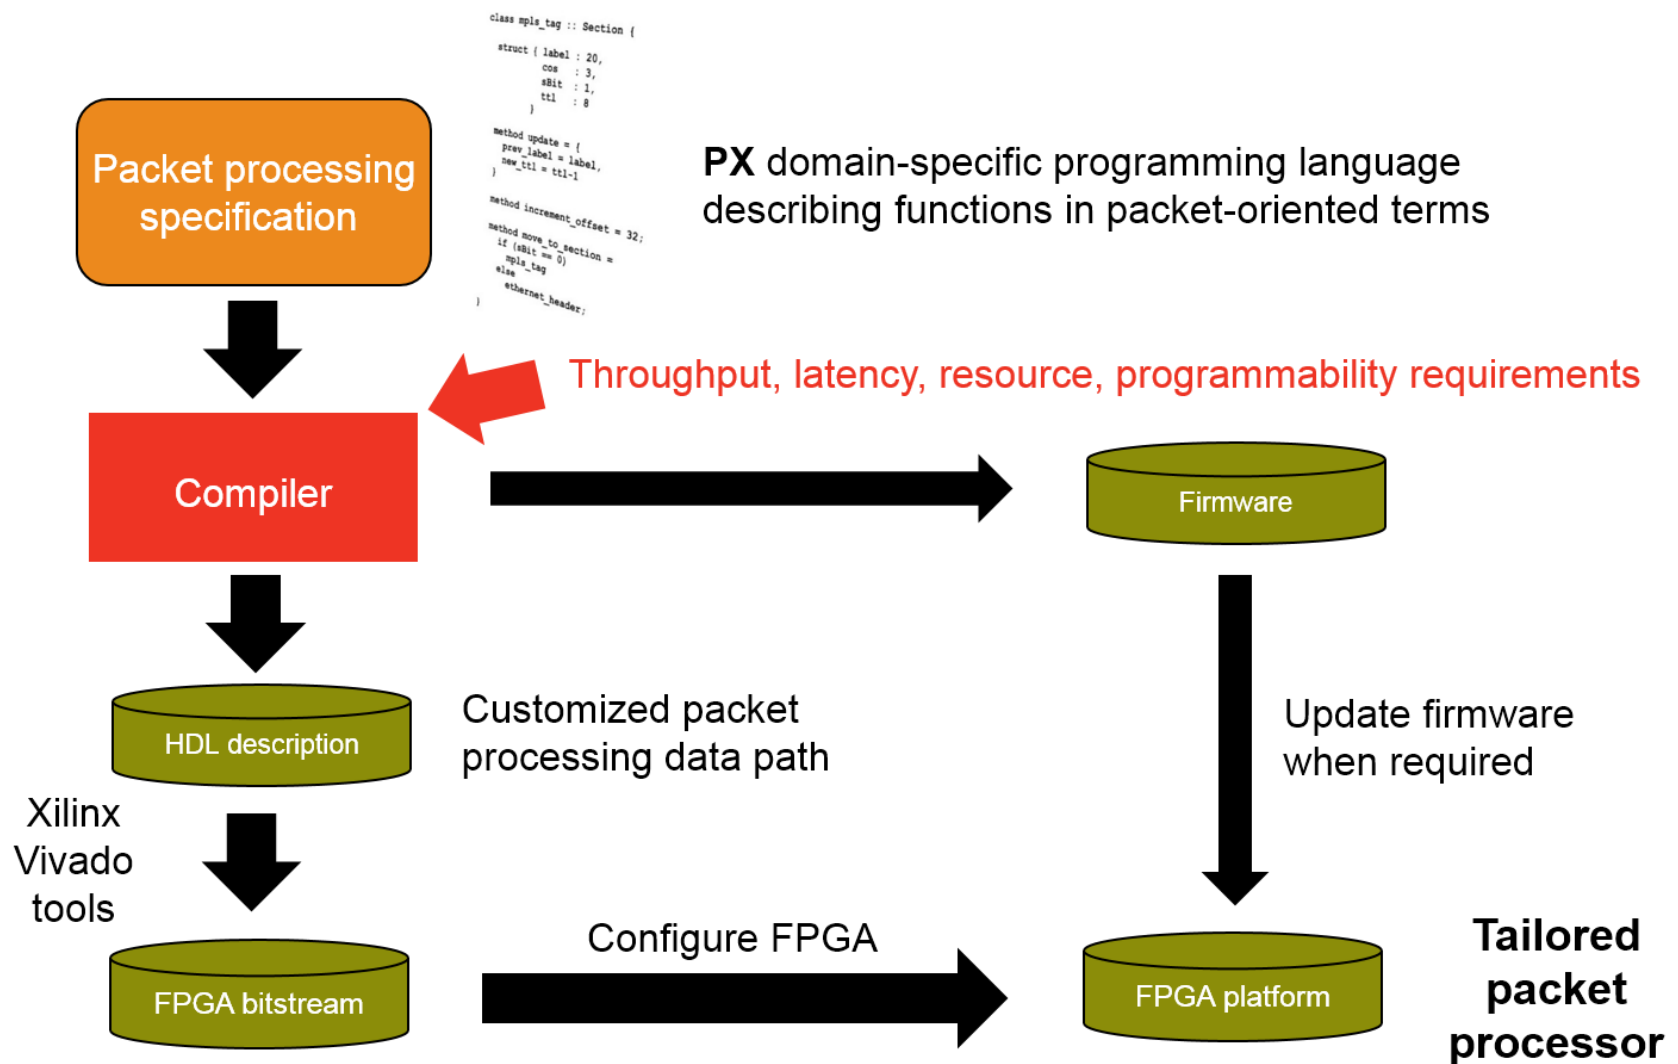
Exact-fit hardware for
reduced cost and power



Firmware for run time
programmability



Design flow and use model



A glimpse of PX

```
struct flow_s {                                // OpenFlow 1.0 12-tuple
    port : 3,                                  // Arrival port
    dmac : 48, smac : 48, type : 16,           // Ethernet
    vid : 12, pcp : 3,                         // VLAN
    sa : 32, da : 32, proto : 8, tos : 6,      // IP
    sp : 16, dp : 16                           // TCP or UDP
}

class OF_parser :: ParsingEngine (9216*8, 4, ETH_header) {

    class OF :: Tuple(inout) { struct flow_s; }
    OF fields;


    // Section sub-class for an Ethernet header
    class ETH_header :: Section {
        struct { dmac : 48, smac : 48, type : 16 }
        method update = {
            fields.dmac = dmac, fields.smac = smac, fields.type = type
        }
        method move_to_section =
            if (type == 0x8100) VLAN_header
            else if (type == 0x0800) IP_header else done(1);
    }

    // Section sub-class for a VLAN header
    class VLAN_header :: Section {
        struct { pcp : 3, cfi : 1, vid : 12, tpid : 16 }
        method update = { fields.vid = vid, fields.pcp = pcp }
        method move_to_section = if (tpid == 0x0800) IP_header else done(1);
    }

    // Section sub-class for an IP header ...
}
```

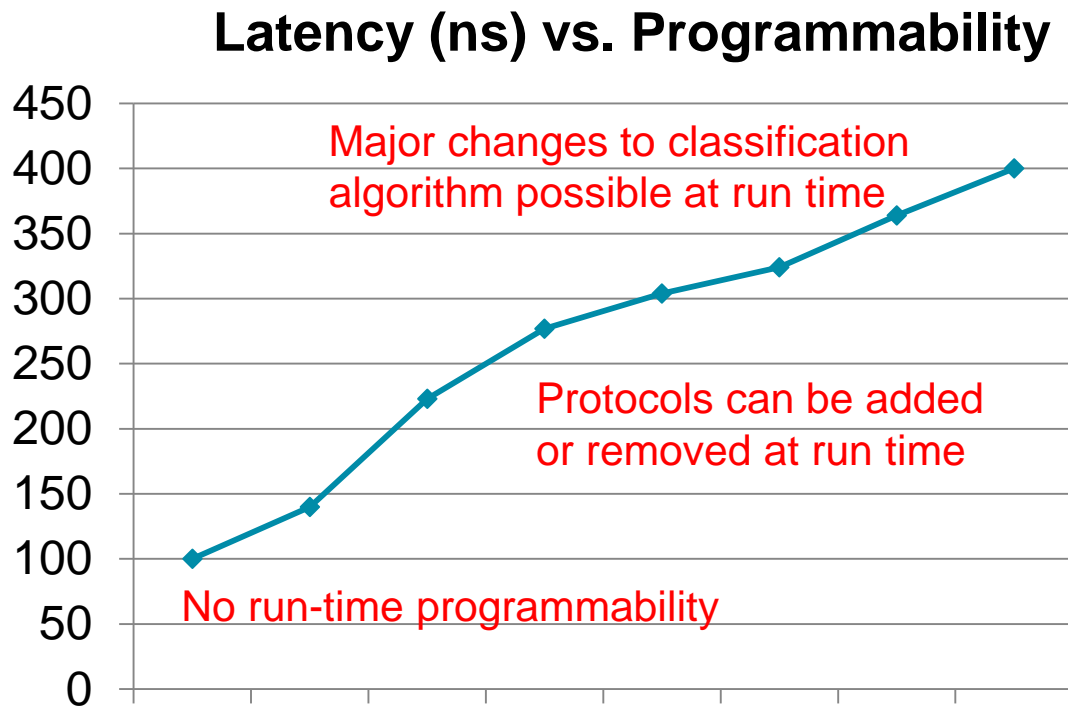

Throughput can be traded off with FPGA utilization

➤ Different raw throughputs through varying data path width parameter



Latency can be traded off with programmability

- High-level description specifies the packet processing functions
- Compiler optimizes implementation characteristics by providing just the required run-time programmability



Example of trade-offs for 30-protocol packet classification design

Current P4 and PX: Many similarities

- Domain-specific: fast path packet processing
 - Declarative: based on packet-centric abstractions
 - Describe per-packet processing in isolation
 - Restricted computation model: no loops, pointers, etc.
 - Abstracted from physical implementation detail
 - *Not designed by language experts*
-
- Headers (“sections” in PX) specified using structs
 - Packet metadata specified using structs
 - Protocol-agnostic packet parsing via programmed transitions
 - Exact match, LPM, and ternary style matching
 - Protocol-agnostic packet editing

Current P4 and PX: Some differences

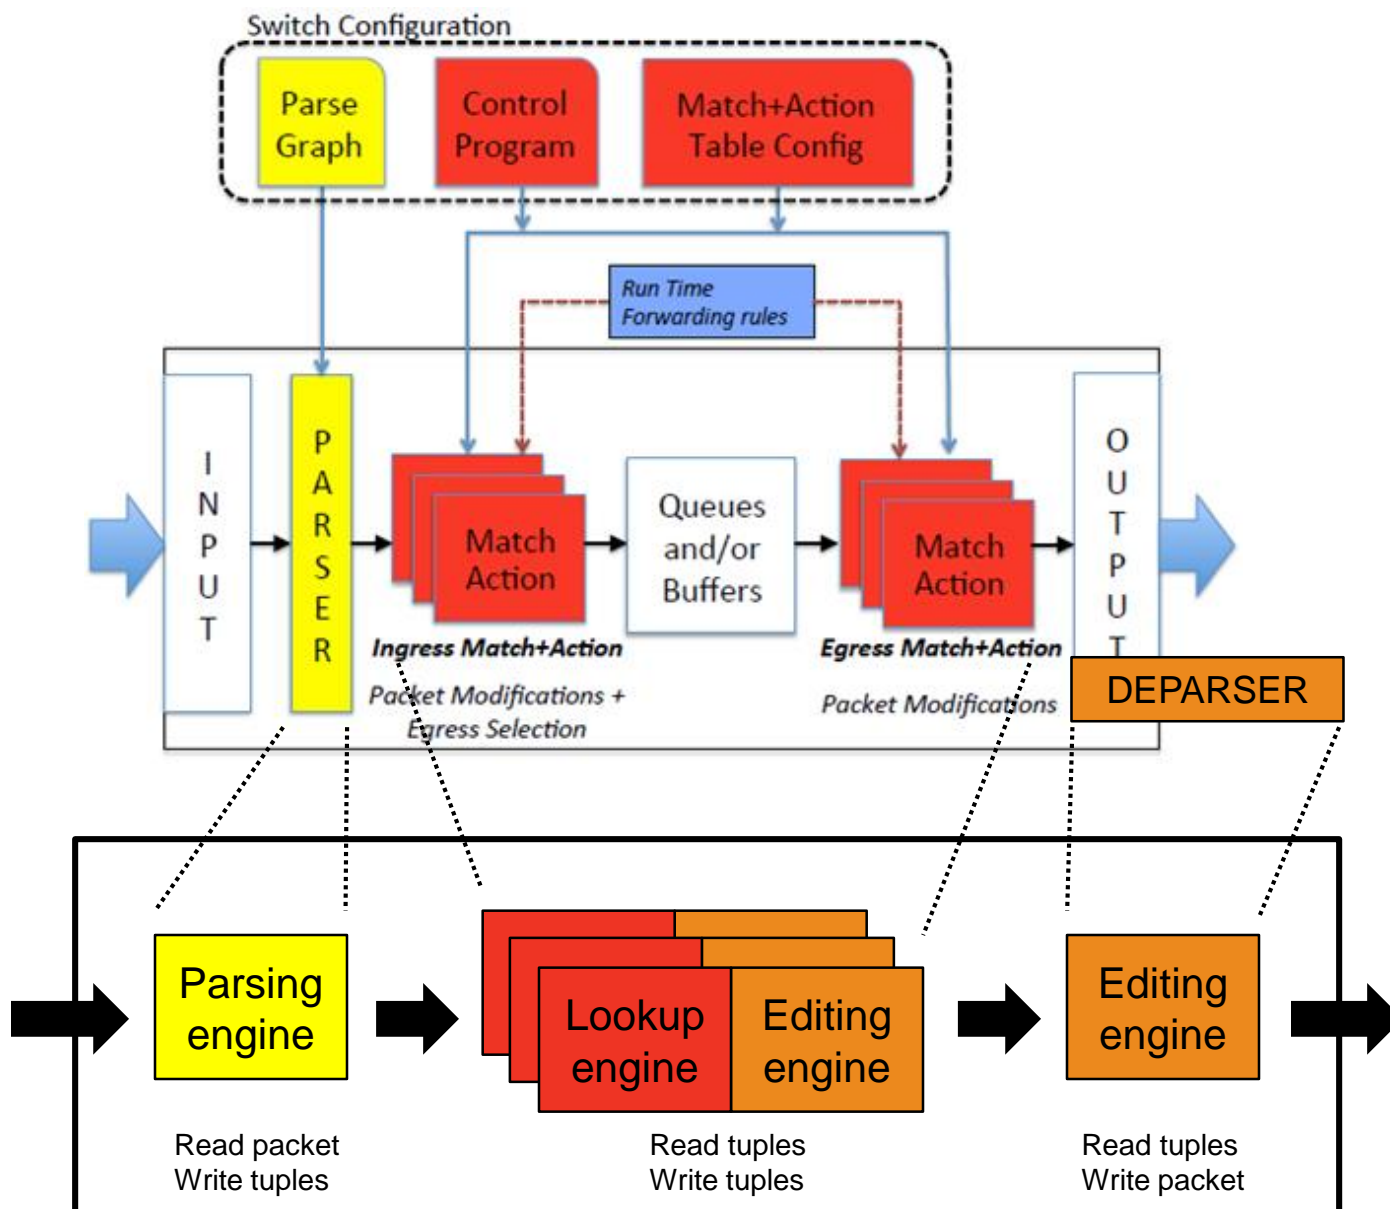
➤ PX features:

- Object-oriented semantics, and modular specifications
- More general parsing transitions, and programmable offset calculation
- Allows incorporation of autonomous (library or user) black box engines
- Hierarchical data flow-based composition of engines to build data planes
- Packet update model based on editing operations (not deparsing)

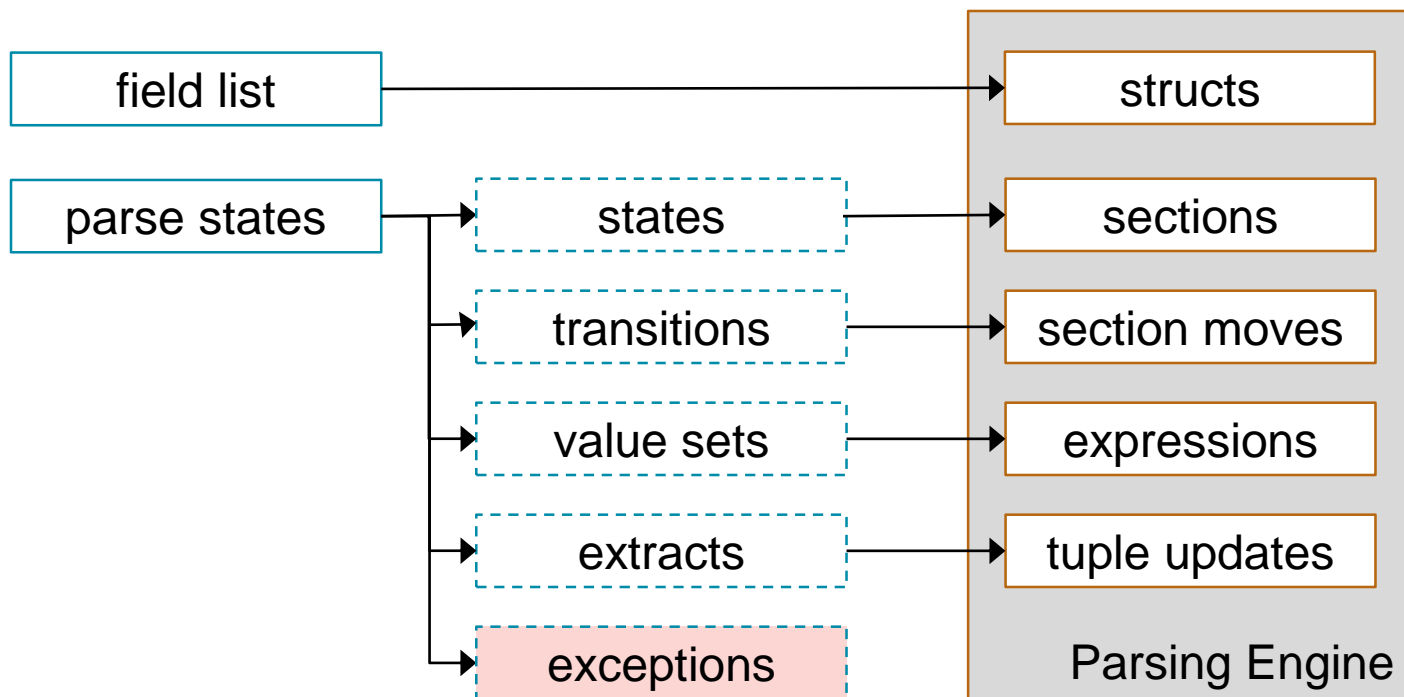
➤ P4 features:

- Match-action table semantics
- Imperative control flow functions
- Built-in complex calculations (e.g., checksums, hashes)
- Explicit inter-packet state: counters, meters, registers
- Resubmit, recirculate, and clone, actions

Mapping P4 to PX

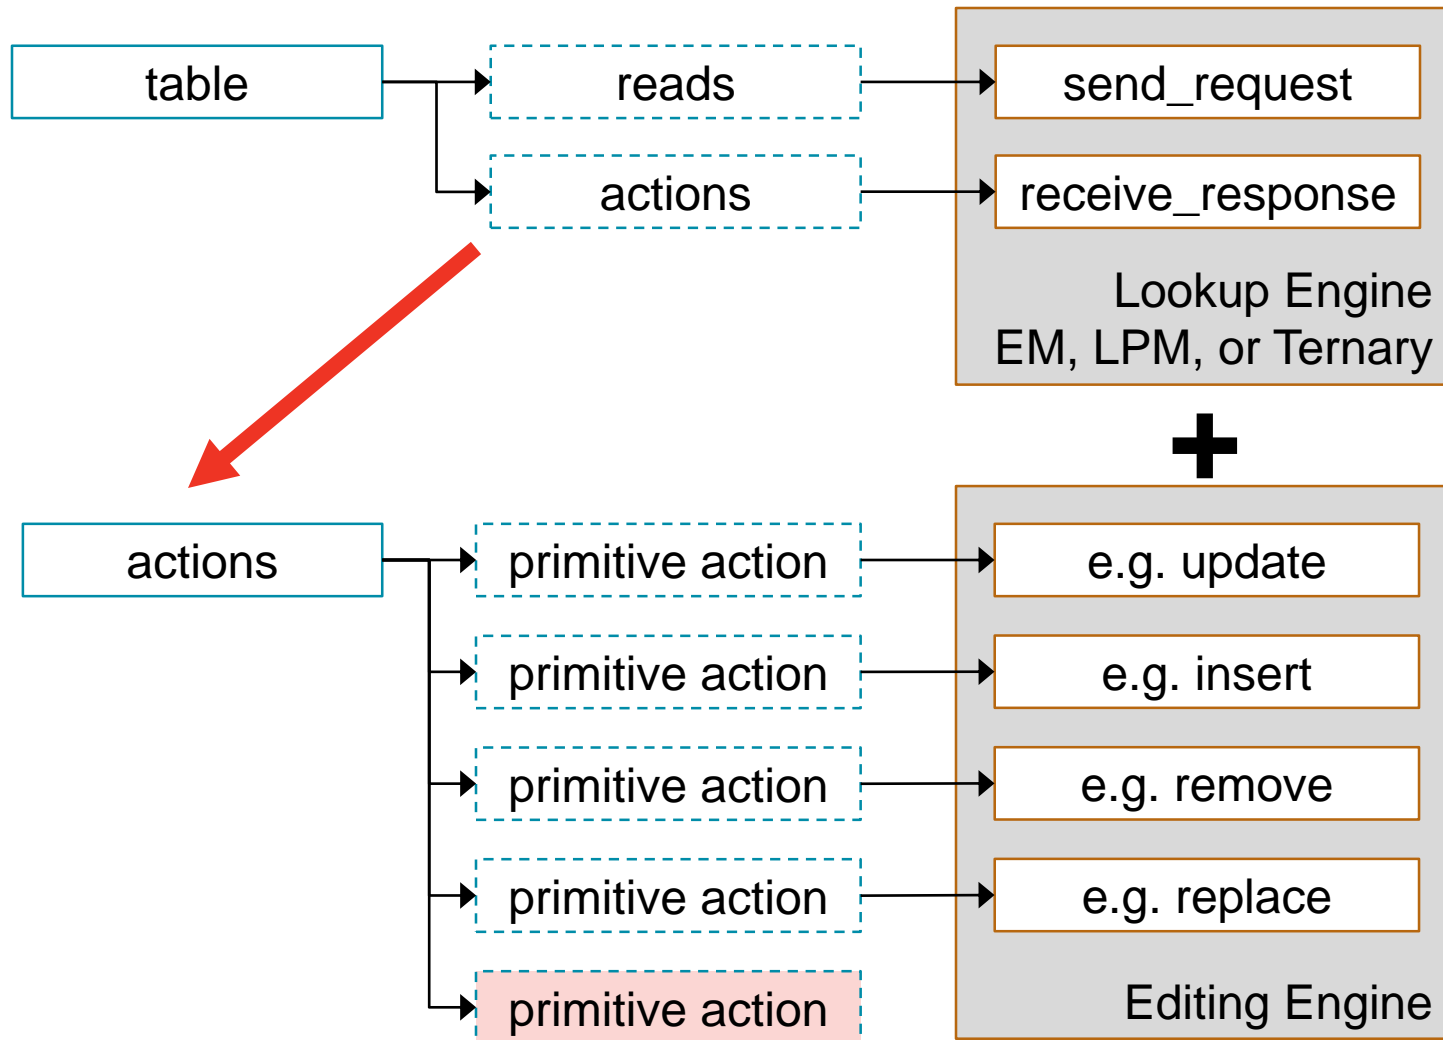


Parser

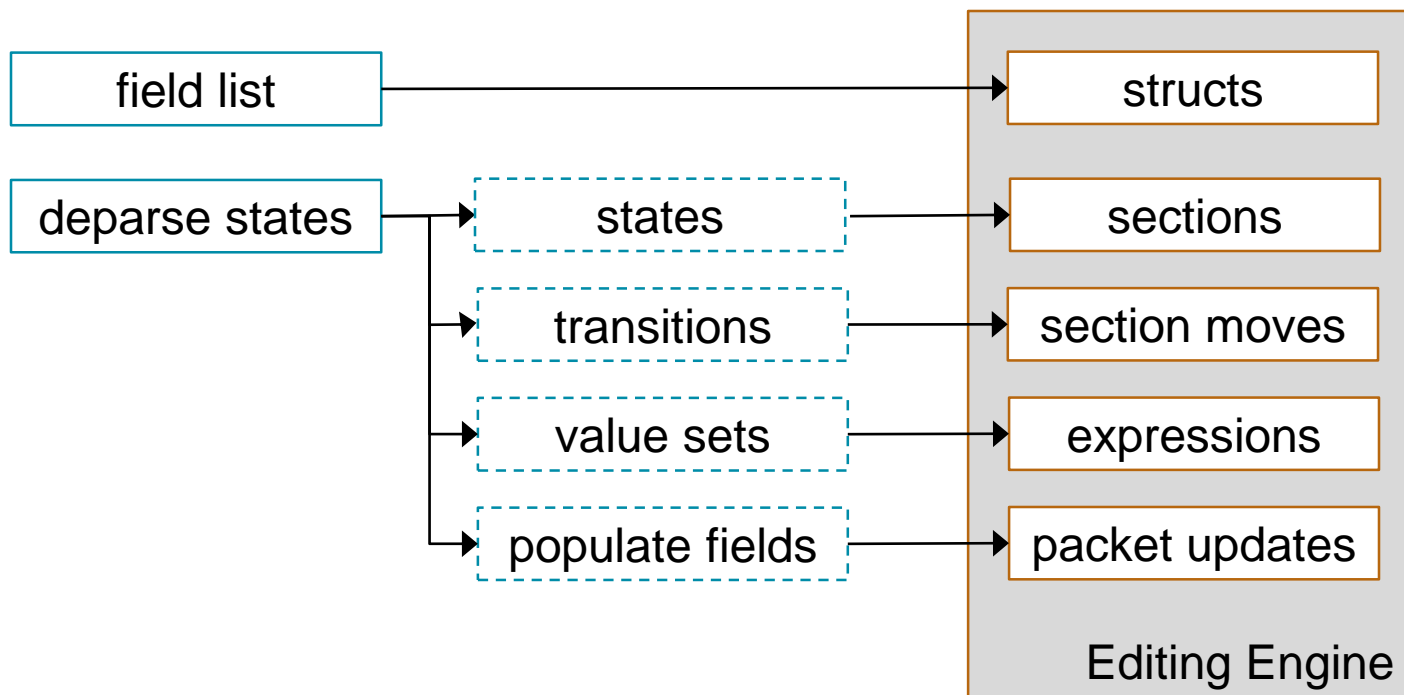


Match-Action table

Mapped to Lookup Engine and Editing Engine pair



Deparser



Optimization experiment 1

➤ Direct mapping to PX:

- Each MAT mapped to a LE-EE pair
- Ternary LEs used for generality

➤ Optimized mapping to PX:

- Small MAT functions included within parsing states (“POF-style”)
- Lookup uses in-expression CAM plus operator-based wildcarding
- Actions coded by tuple updates
- Deparser editing engine does all packet update actions

➤ Example impact: 23% resource reduction

Optimization experiment 2

➤ Direct mapping to PX:

- Parser extracts all fields, and Deparser reconstructs all fields
 - All fields passed along data path
- MATs have deferred actions on packets:
 - Either direct updates to fields being sent along data path ...
 - ... or coded deferred actions passed down data path

➤ Optimized mapping to PX:

- Parser extracts only fields needed for MAT keys, or other expressions
 - Packets passed along data path
- MATs have immediate actions on packets:
 - Either direct updates to packets being sent along data path ...
 - ... or coded 'end of processing' actions passed down data path

➤ Example impact: 41% resource reduction

Prototype demonstrated today



P4.org front end



HILIR



Xilinx Labs mapper



➤ Front end: github.com/p4lang/p4-hilir


➤ Mapper: new code written in Python

➤ SDNet: next product release version

➤ Xilinx VCU109 development board:

- Carries Virtex Ultrascale XCVU095 FPGA
- Has quad CFP4 interface
 - Being used for 400G Ethernet demos

Xilinx
SDNet



Prototype demonstrated today

➤ Normally, connect board to Ixia tester via single CFP4:

- Tester generates input minimum-size packets at 100G line rate
- Packets pass through fast path generated from P4 description
- Tester captures output packets at 100G line rate

➤ Today, the compact and quiet version:

- Test packets written to FPGA
- On-FPGA block generates input minimum-size packets at 100G line rate
- Packets pass through fast path generated from P4 description
- On-FPGA block captures and stores output packets at 100G line rate
- Captured packets read from FPGA
- Wireshark shows packets before and after processing

➤ MicroBlaze soft processor on FPGA is the on-chip controller

Protocol Independent Forwarding (PIF)

➤ “How to tell your plumbing what to do”

- Nick McKeown, ONF member workday keynote, September 2014

➤ Nick’s observation:


- Some switches are more programmable than fixed-function ASICs

➤ Proposed PIF approach:

- **Phase 0.** Initially, the switch does not know what a protocol is, or how to process packets (Protocol Independence)
- **Phase 1.** We tell the switch how we want it to process packets (Configuration)
- **Phase 2.** The switch runs (Run-time)

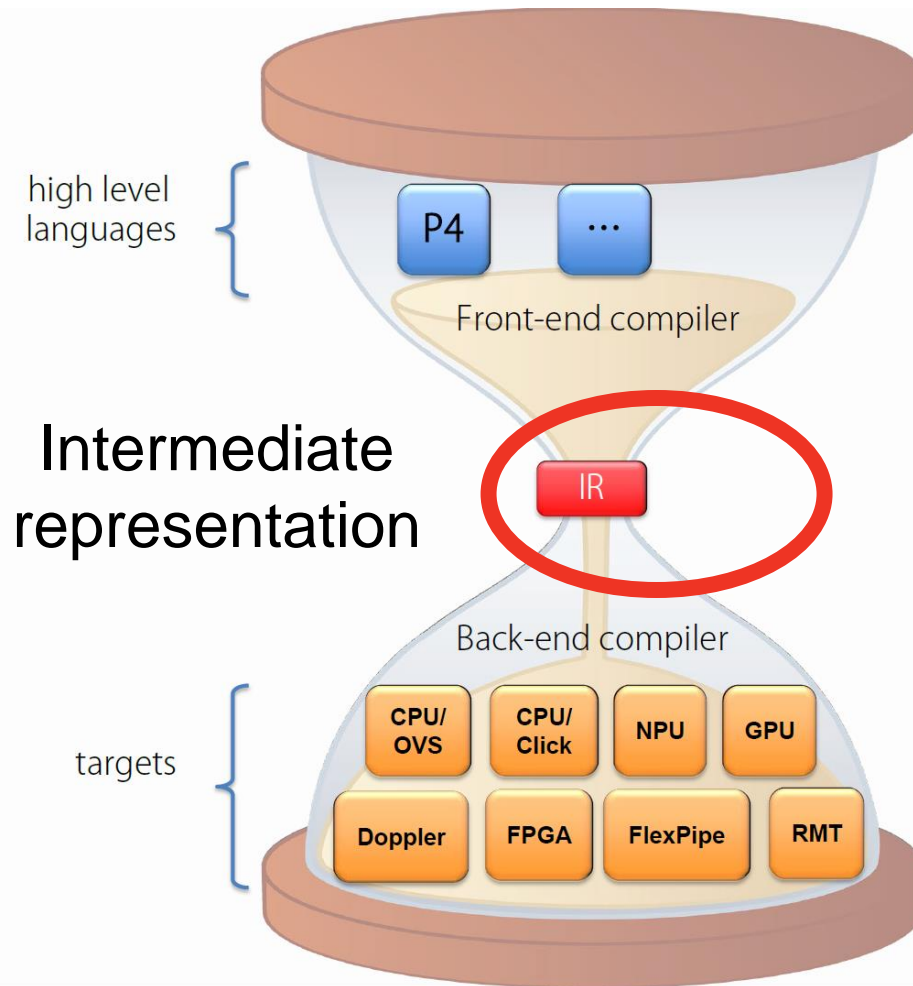
PIF open source software project

Within ONF Oct 2014 to Apr 2015, open to all since May 2015



OPEN SOURCE SDN

opensourcesdn.org



Focus areas of PIF project:

- Experimenting with IR elements
- Use cases
- Runtime APIs




Feeds into ONF
“OpenFlow Next Gen”
specification activity

P4-to-Xilinx FPGA direct compilation flow



Open source
front end



➤ Moving target ...

➤ Anticipated ...

➤ Moving target ...

➤ Refactoring ...

➤ Range of Xilinx FPGAs

Summary

- **FPGA is compelling target technology**
 - Combination of programmability with performance
- **Demonstration of P4-to-Xilinx flow generating 100G fast path**
 - Prototype leveraging existing Xilinx SDNet methodology
- **Active participation in P4 evolution**
- **Active participation in PIF project**
- **Xilinx invites researchers to undertake collaborative projects**

