

P4GPU: A Study of Mapping a P4 Program onto GPU Target

Peilong Li, Tyler Alterio, Swaroop Thool and Yan Luo

ACANETS Lab (<http://acanets.uml.edu/>)

University of Massachusetts Lowell



Background & Motivation

- P4 enables programmability of network data plane
- Software-based routing and forwarding:
 - Programmability ↑
 - Line rate ↓
- P4 hardware targets:
 - CPU, NPU, reconfigurable switches, iNICs, FPGA, etc.
 - GPGPU potentiality? – Hundreds/thousands of cores, SIMD architecture, high memory bandwidth, mature parallel programming environment.

Contributions

- ***P4-to-GPU***: provide a viable way to map a portion of the P4 code onto GPU target.
- ***Acceleration***: design a heterogeneous framework with CPU and GPU to accelerate performance.
- ***Optimization***: optimize lookup and classification kernels on GPU and explore load balancing on the CPU/GPU framework.

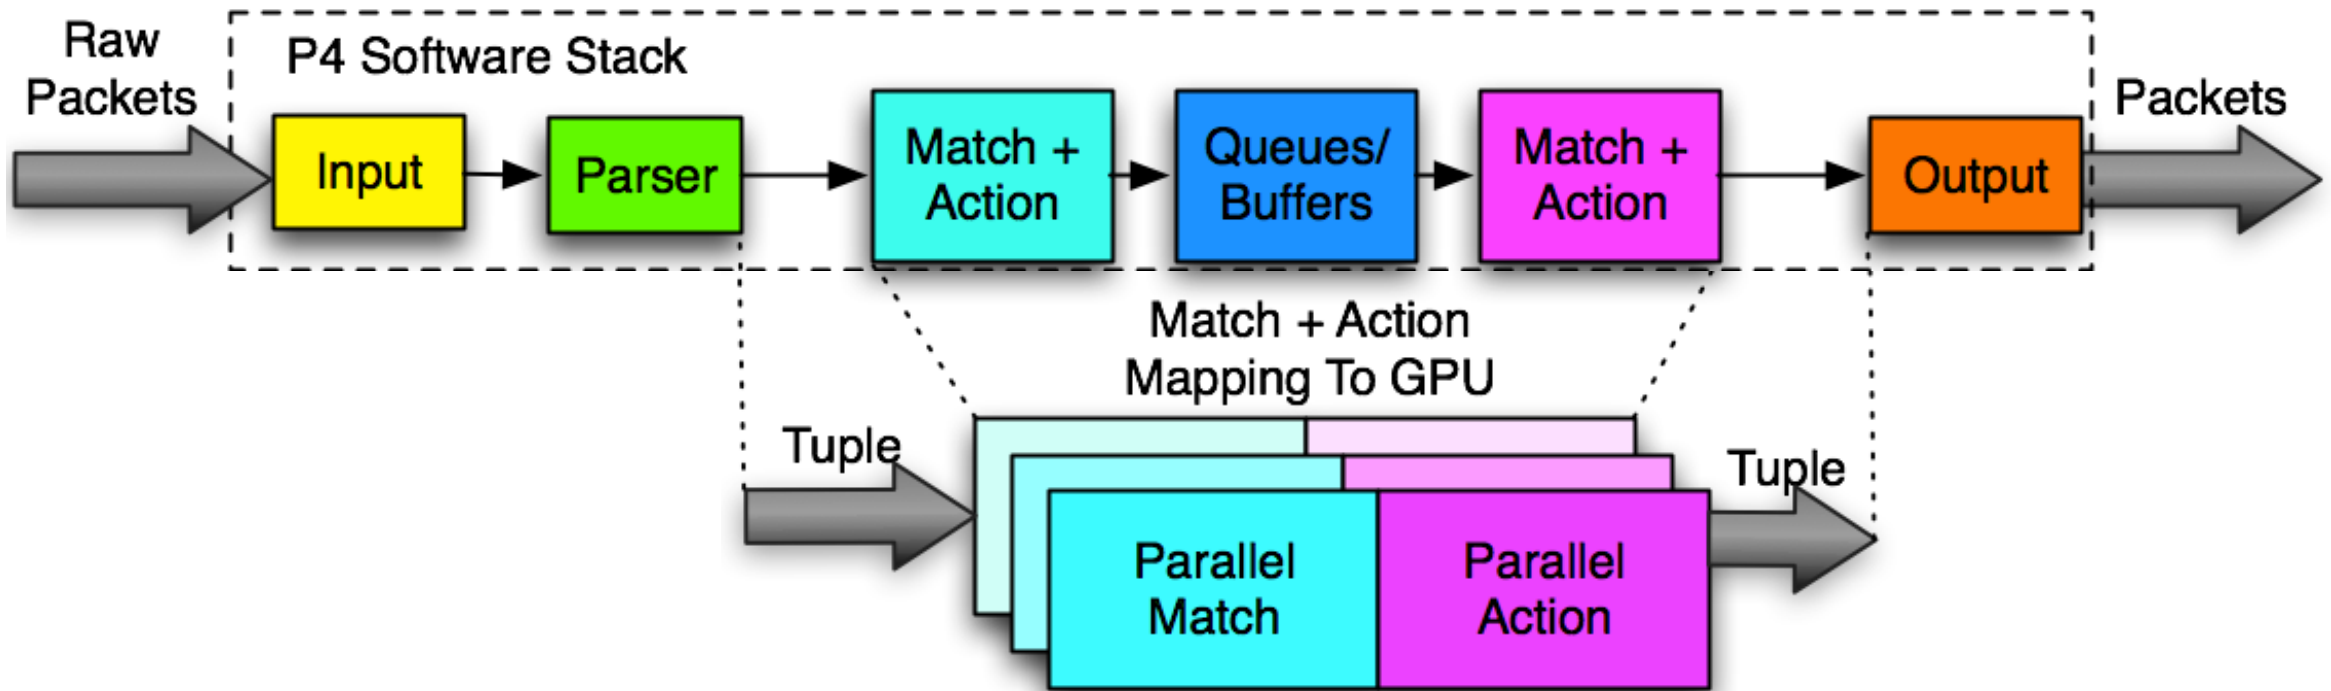
P4 to GPU: Overview

- Mapping P4 to GPU difficulties & hints

Difficulties	Hints
Lack of dedicated hardware (match+action)	Take advantage of many cores
Inefficient kernel design leads to performance disaster	Kernel requires optimization; auto-generated kernel may not render good performance
Coprocessor has limited memory and low clock frequency	Heterogeneous CPU/GPU architecture is desired
GPU is good at high parallelism (table lookup, matching ...); bad at dependency and branching (parsing, state machine ...)	Only offload parallelizable code to GPU.

P4 to GPU: Software Stack

- Partially mapping




P4 to GPU: Design

- Three steps of mapping:


Step 1: P4 intermediate representation preparation

Step 2: GPU kernel configuration generation

Step 3: GPU kernel initialization



P4 to GPU: Code Diagram






- The GPU kernel conf parser:
 - Parse h.p4_tables etc. in IR to obtain P4 table conf
 - Find the order of tables * from the control flow.

* Sequential tables for now


P4 to GPU: GPU Kernel Design

- LPM lookup kernel:
 - Used for “*ipv4_lpm*” table
 - Designed for both IPv4 and IPv6
 - Baseline kernel: linear search
 - Optimized kernel: binary trie and k-stride multibit trie
- Rule based classifier kernel:
 - Used for “*forward*” and “*send_frame*” table
 - Designed for both exact and wildcard match
 - Baseline kernel: linear search
 - Optimized kernel: grid-of-tries



P4 to GPU: GPU Kernel Design


- Latency hiding:
 - Batch processing: reduce number of data copying
 - 2D pipelining:



- Memory tweaking:
 - Use texture memory for storing lookup tables
 - Use mapped memory from host to device to reduce data movement overhead
- Data structure:
 - Hash tables/tries are not GPU-favored data structure
 - Trie-to-array vectorization

Architectural Design

- Components: CPU, main memory, GPU




Step 1: receiving
Step 2: batching
Step 3: load-balancing and offloading
Step 4: results buffering
Step 5: forwarding actions

Evaluation: Setup


- Experiment platform:
 - CPU: Intel Quad Core i7-3610 QM @ 2.3 GHz
 - Main memory: 8 GB
 - GPU: Nvidia GT 650M with 384 cores
- Dataset:
 - RouteView (Jan 2015) IPv4 (550,000 entries), IPv6 (20,000 entries)
 - ClassBench: a set of filters, i.e. FW and ACL
- Traffic generation:
 - Random generator: assumes no packet IO overhead
 - Click Modular Router: emulated packet sending/receiving (socket-based)

Evaluation: Results

- Lookup Kernel





- Classifier Kernel



Evaluation: Results

- Throughput and Latency



Conclusion & Future Work

- P4GPU tool provides a viable way to map P4 code onto GPU target.
- GPU is a promising target for P4 to achieve high throughput and low latency forwarding.
- Integrate P4 table dependency to P4GPU tool
- Study the possibility of optimized GPU kernel auto generation.
- Implement GPU kernels for more functionalities, e.g. OpenFlow.

That's All

