



XILINX

ALL PROGRAMMABLE™

Programmable Target Architectures for P4

Gordon Brebner
Xilinx Labs
San José, USA

P4 Workshop, Stanford University, 18 November 2015

What this talk is about

- **Recap: P4-to-FPGA compilation demonstrated at June workshop**
- **Architecture/language decoupling discussion**
- **Click, and extended port/connection semantics**
- **Xilinx SDNet's PX data plane building**
- **Prototype 1: Generation of P4 includes from PX description**
- **Prototype 2: Generation of PX from Extended P4 description**

Prototype demonstrated at June P4 Workshop



P4.org front end



HILIR




Xilinx Labs mapper

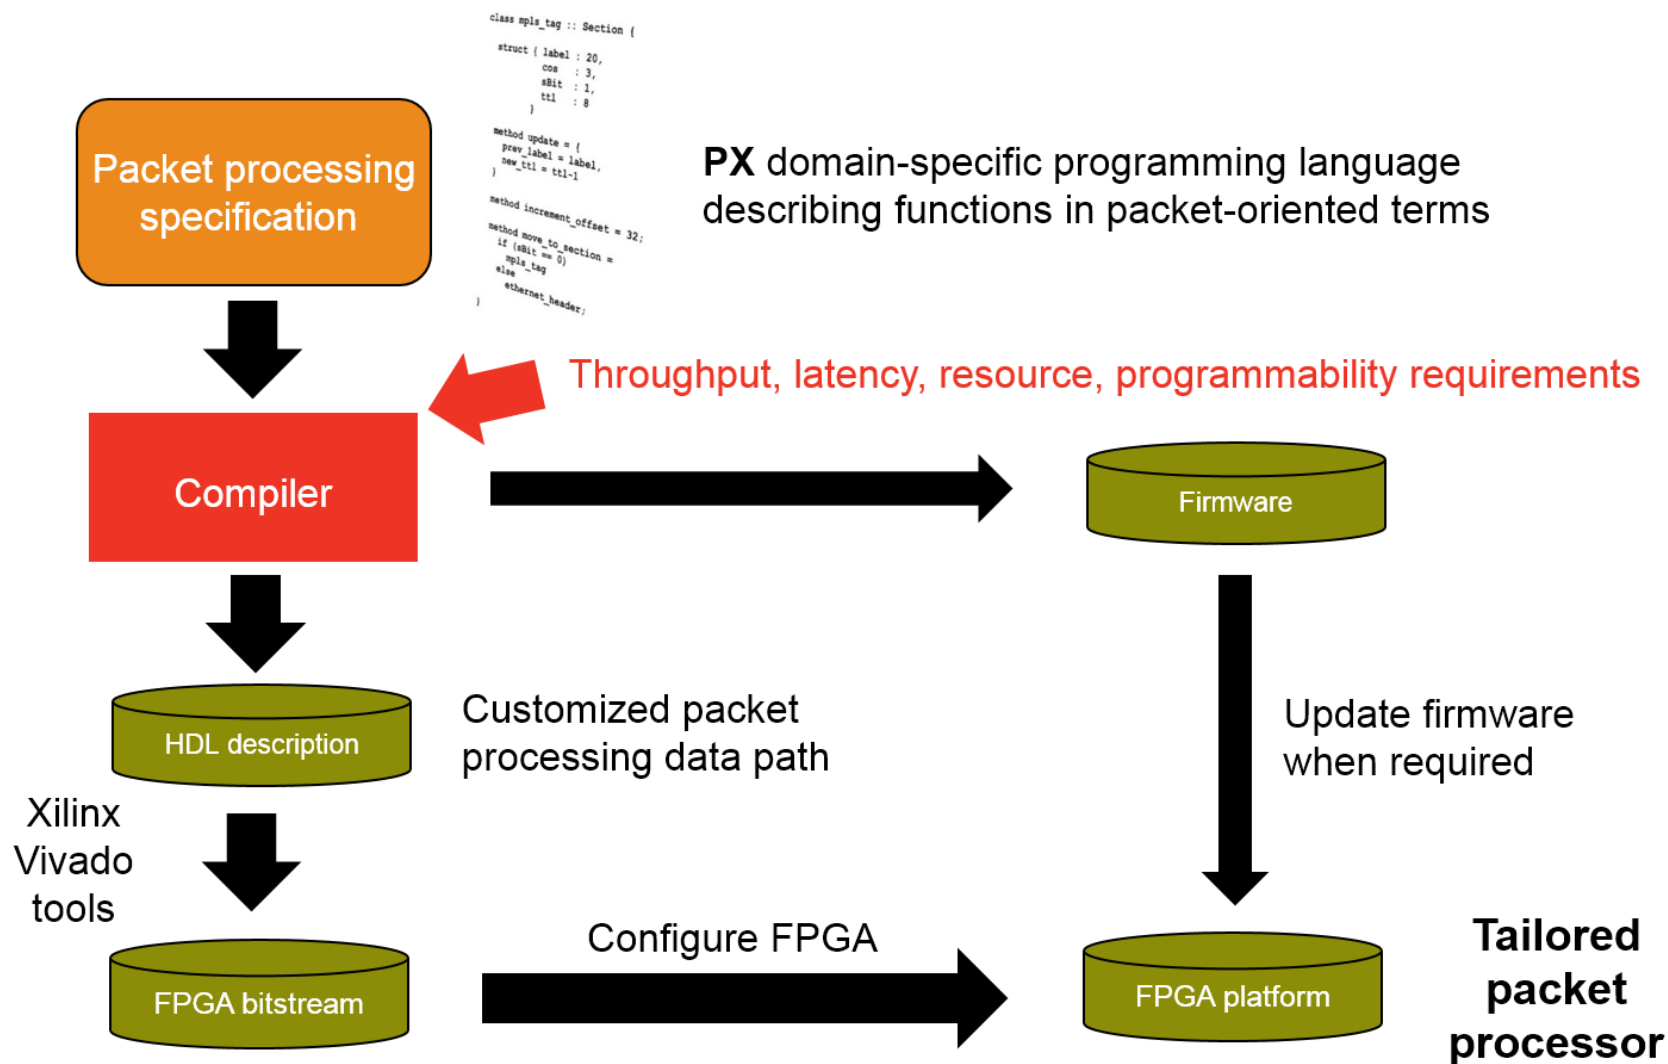
Xilinx
SDNet



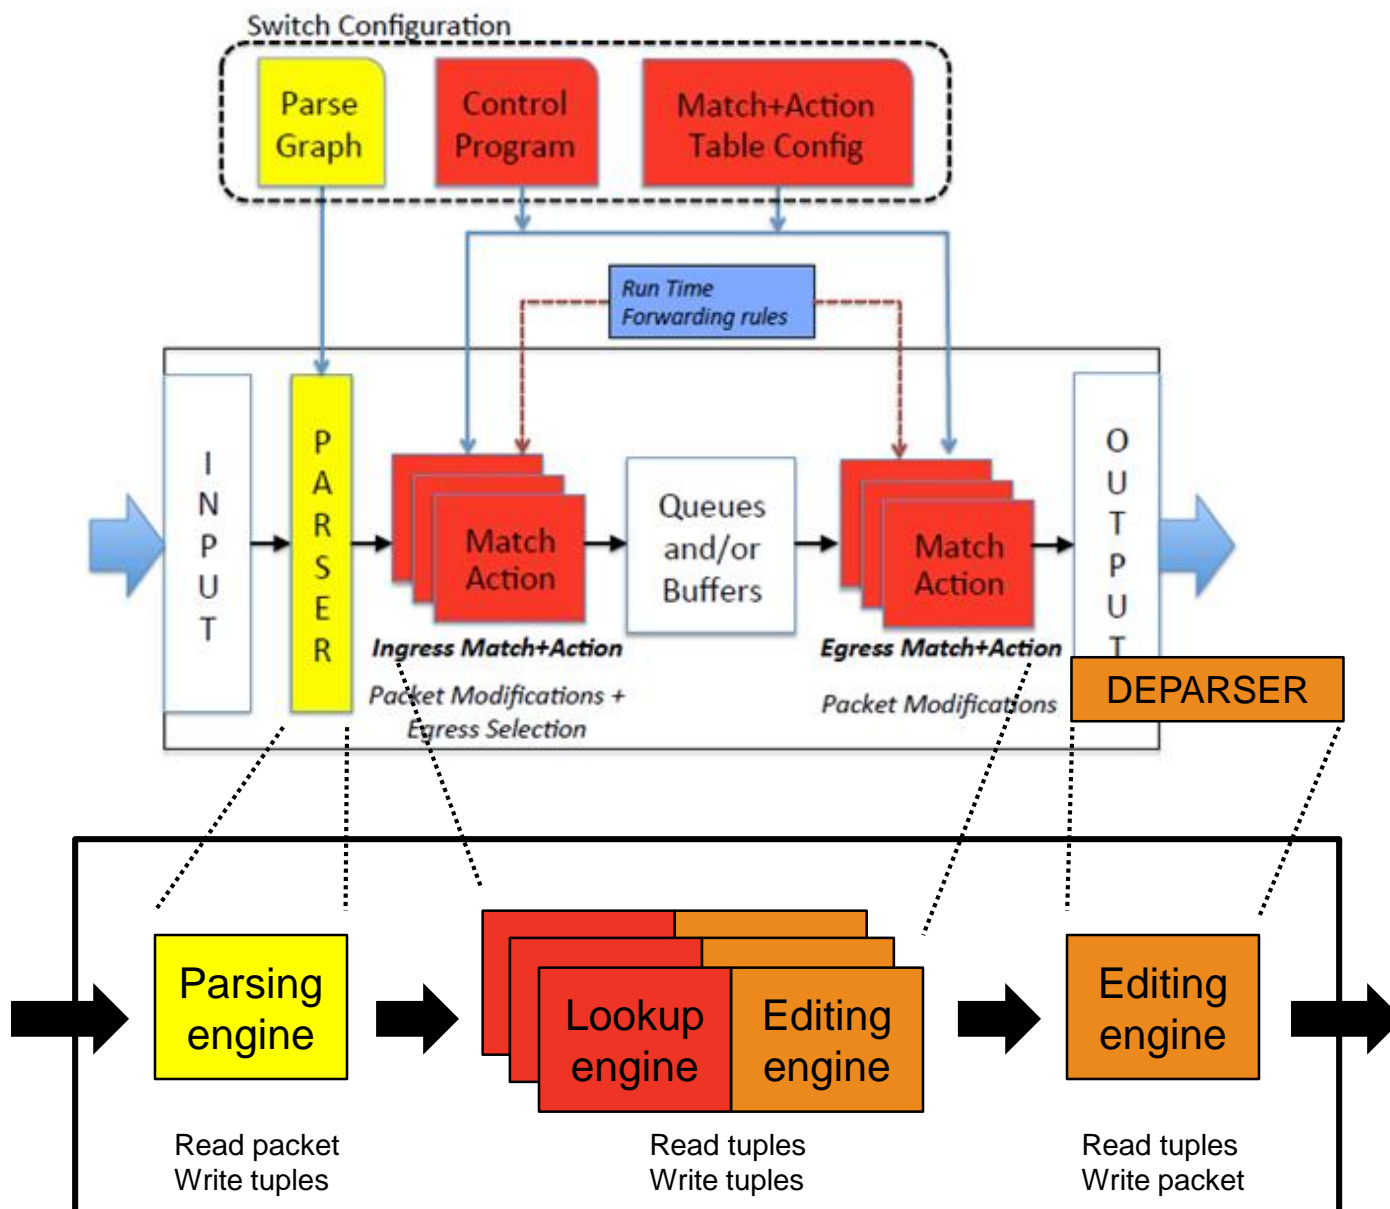
- Front end: github.com/p4lang/p4-hlir
- Mapper: new code written in Python
- SDNet: next product release version
- Xilinx VCU109 development board:
 - Carries Virtex Ultrascale XCVU095 FPGA
 - Has quad CFP4 interface



Xilinx SDNet design flow and use model



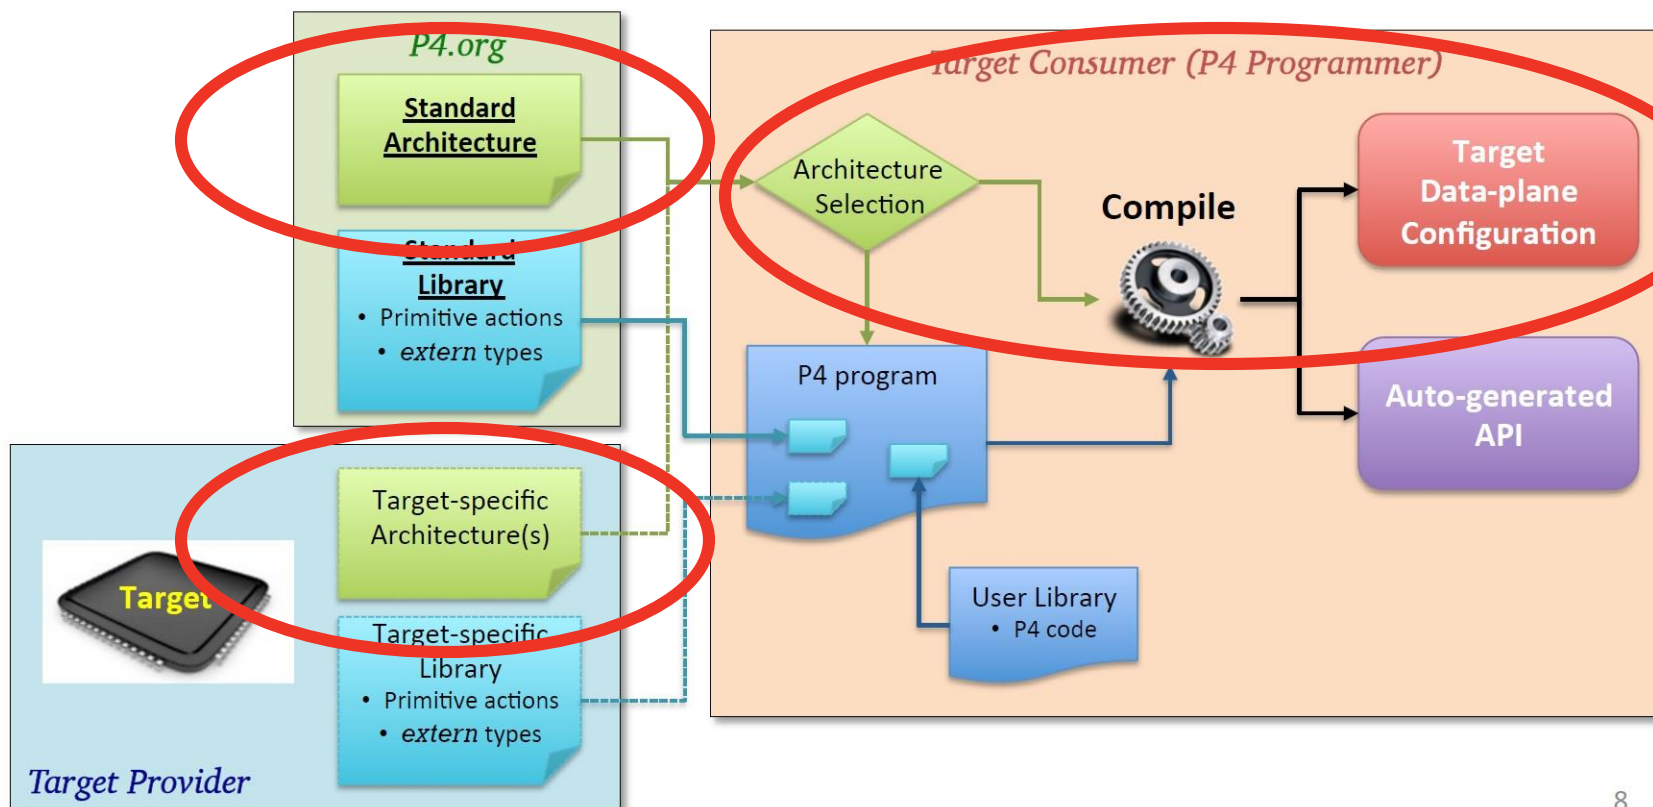
Mapping P4 to PX



Proposed P4 architecture/language separation

➤ Slide from Chang's earlier talk today:

Fitting all these together



Standard Click, and describing architectures

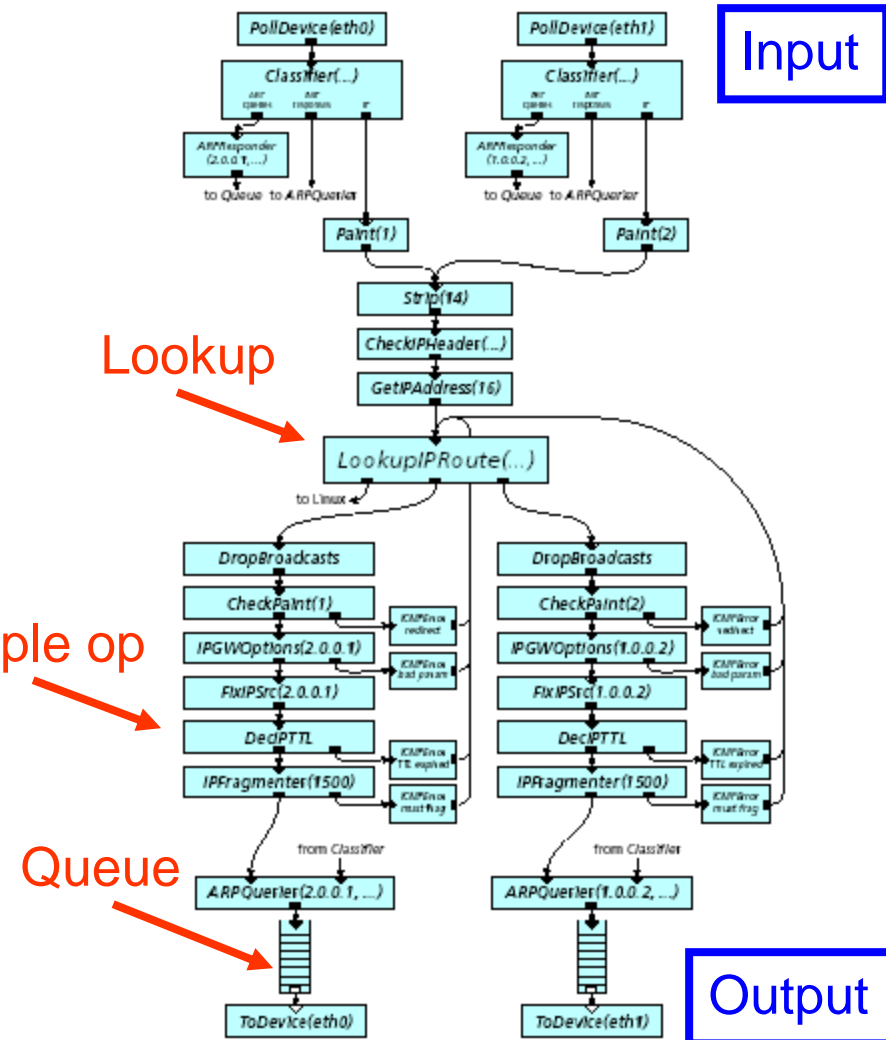
Each box is an instance of a pre-defined Click *element*

Connections between elements, and packets 'pushed' and 'pulled' through the resulting graph

Some adverse characteristics:

- Fixed-function elements
- Fine-grain element functions
- Inter-element interaction:
 - data flow for packets
 - *method calls otherwise*

Click-to-FPGA research (DAC 2004)



Extended Click connection semantics

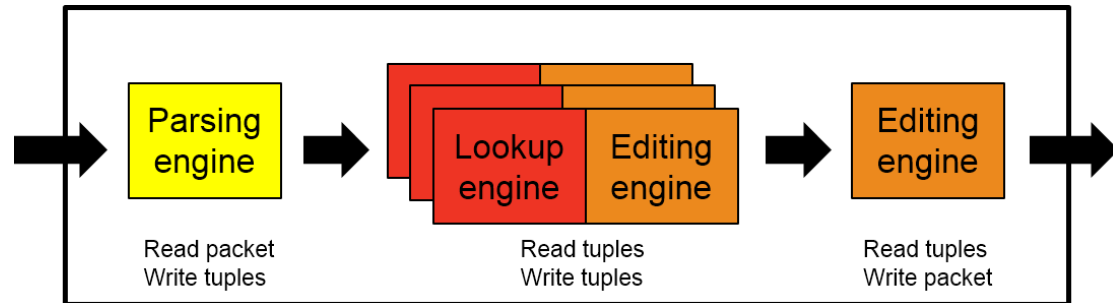
- **Xilinx Labs work (FCCM 2010, ACM TRETs 2012)**
- **Five types of connection between ports:**
 - Packet (as in Standard Click)
 - Tuple (as used for metadata in P4, for example)
 - Access (store access: read and write operations)
 - Compute (procedure call and return)
 - Plain (no semantics, just wiring – low-level escape mechanism)
- **Featured with Click syntax in Xilinx *Packet Xpress* (FPT 2009)**
- **Features with PX syntax in Xilinx SDNet**
 - Only packet and tuple types available to user

Example: Data plane building in PX

(Auto-generated from P4 1.0 description by P4-to-PX mapper)

```
class p4_processor :: System {
  Packet_input instream;
  Packet_output outstream;

  Parser parsing_engine1;
  Deparser editing_engine2;
  update_eth_dest lookup_engine3;
  process_ipv4 editing_engine6;
  next_hop_ipv4 lookup_engine5;
  process_eth editing_engine4;
```



```
method connect = {
  parsing_engine1.packet_in = instream,
  outstream = editing_engine2.packet_out,
  editing_engine4.wire2_tupleIn = lookup_engine3.response,
  lookup_engine3.request = parsing_engine1.update_eth_dest_tupleOut,
  editing_engine6.wire1_tupleIn = lookup_engine5.response,
  lookup_engine5.request = parsing_engine1.next_hop_ipv4_tupleOut,
  editing_engine4.standard_metadata_tupleInOut = parsing_engine1.standard_metadata_tupleOut,
  editing_engine4.ethernet_tupleInOut = parsing_engine1.ethernet_tupleOut,
  editing_engine4.ipv4_tupleInOut = parsing_engine1.ipv4_tupleOut,
  editing_engine4.packet_in = parsing_engine1.packet_out,
  editing_engine6.standard_metadata_tupleInOut = editing_engine4.standard_metadata_tupleInOut,
  editing_engine6.ethernet_tupleInOut = editing_engine4.ethernet_tupleInOut,
  editing_engine6.ipv4_tupleInOut = editing_engine4.ipv4_tupleInOut,
  editing_engine6.packet_in = editing_engine4.packet_out,
  editing_engine2.standard_metadata_tupleIn = editing_engine6.standard_metadata_tupleInOut,
  editing_engine2.ethernet_tupleIn = editing_engine6.ethernet_tupleInOut,
  editing_engine2.ipv4_tupleIn = editing_engine6.ipv4_tupleInOut,
  editing_engine2.packet_in = editing_engine6.packet_out
}
```

Prototype 1: Generate P4 includes from PX (1)

➤ Mark certain PX components with P4 significance

➤ P4 engine: a component described in P4

```
class Parser :: P4Engine (parser) {  
    StandardOut standard;  
    EthernetOut ethernet;  
    IPv4Out ipv4;  
  
    method classify = { standard=metadata, ethernet=header, ipv4=header }  
}
```

➤ P4 procedure: an extern made available to P4 components

```
class IPv4Checksum :: P4Procedure {  
    IPv4InOut ipv4;  
  
    ChecksumEngine checksum;  
  
    method apply = {  
        checksum.ipv4 = ipv4  
    }  
}
```

Prototype 1: Generate P4 includes from PX (2)

➤ Target-specific P4 includes generated by PX compiler

➤ Header_type declarations (e.g. Standard_t)

➤ From P4 engine (P4-discuss syntax):

```
whitebox_type Parser (  
    out metadata Standard_t standard,  
    out header Ethernet_t ethernet,  
    out header IPv4_t ipv4  
);
```

➤ From P4 procedure (P4 1.1 syntax):

```
extern_type IPv4Checksum {  
    method apply(IPv4_t ipv4);  
}
```

Prototype 2: Generate PX from Extended P4 (1)

➤ Target architecture described in P4 “1.arch”:

- Header type declarations (standard P4 1.1 syntax)
- *Element* type declarations (Click-inspired: target architecture components)
 - Note: what was generated from PX in Prototype 1
- Extern type declarations (standard P4 1.1 syntax)
 - Note: what was generated from PX in Prototype 1
- *Architecture* type declaration (Click-inspired: overall target architecture)
 - Input and output ports
 - User-provided *intern* elements
 - Target-supplied elements
 - Connections between elements
 - Target-provided *extern* objects

Prototype 2: Generate PX from Extended P4 (2)

➤ Example element type declarations

```
element_type Parser (  
    in packet packet_in,  
    out metadata standard_metadata_t standard,  
    out header ethernet_t ethernet,  
    out header ipv4_t ipv4  
);
```

```
element_type Processor (  
    inout metadata standard_metadata_t standard,  
    inout header ethernet_t ethernet,  
    inout header ipv4_t ipv4  
);
```

Prototype 2: Generate PX from Extended P4 (3)

➤ Example architecture type declaration

```
architecture_type IP_pipeline {  
    // External interfaces to the architecture  
    in packet instream;  
    out packet outstream;  
  
    // Two user-provided elements  
    intern Parser parser;  
    intern Processor processor;  
  
    // One supplied element  
    Deparser deparser;  
  
    // Connectivity of the architecture  
    connections {  
        (instream, parser.packet_in),           // Input to parser  
        (parser.standard, processor.standard),  // \  
        (parser.ethernet, processor.ethernet),  // Parser to processor  
        (parser.ipv4, processor.ipv4),          // /  
        (processor.standard, deparser.standard), // \  
        (processor.ethernet, deparser.ethernet), // Processor to deparser  
        (processor.ipv4, deparser.ipv4),         // /  
        (deparser.packet_out, outstream)        // Deparser to output  
    }  
  
    // One provided extern  
    extern checksum ipv4_checksum;  
}
```

Prototype 2 demonstrated today at 100 Gb/s rate



Extended
HLIR




P4 1.arch

Front end

Xilinx Labs mapper

```
architecture_type IP_pipeline {  
  // External interfaces to the architecture  
  in packet instream;  
  out packet outstream;  
  
  // Two user-provided elements  
  intern Parser parser;  
  intern Processor processor;  
  
  // One supplied element  
  Deparser deparser;  
  
  // Connectivity of the architecture  
  connections {  
    (instream, parser.packet_in),  
    (parser.standard, processor.standard),  
    (parser.ethernet, processor.ethernet),  
    (parser.ipv4, processor.ipv4),  
    (processor.standard, deparser.standard),  
    (processor.ethernet, deparser.ethernet),  
    (processor.ipv4, deparser.ipv4),  
    (deparser.packet_out, outstream)  
  }  
  
  // One provided extern  
  extern IPv4Checksum checksummer;  
}
```

Xilinx
SDNet



Summary

- **P4 heading towards architecture/language separation**
 - As explained by Chang earlier today
- **Strawman proposal for describing architecture using P4 “1.arch”**
 - As opposed to using English – more precise, and machine processible
 - Inspired by the well-known Click element/connection paradigm
- **Demonstration of compiling P4 “1.arch” to FPGA at 100G rate**
 - Including user P4 components written in P4 “1.user” (more-or-less 1.1)
- **Xilinx Labs invites researchers to undertake collaborative projects**

