

# Packet Transactions: Programming the Data Plane at Line Rate

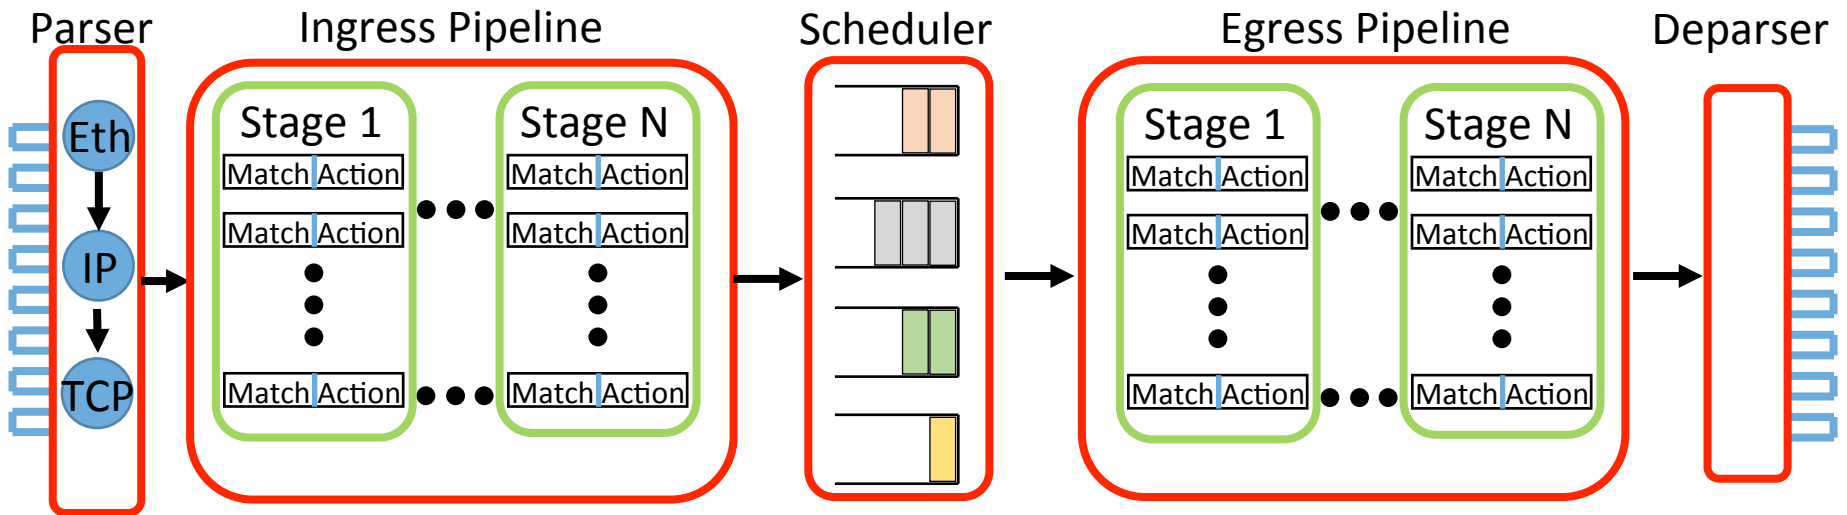
**Anirudh Sivaraman**, Mihai Budiu, Alvin Cheung,  
Changhoon Kim, Steve Licking, George Varghese,  
Hari Balakrishnan, Mohammad Alizadeh, Nick  
McKeown



# Programming the data-plane at line rate

- Programmable: Can we express a new data-plane algorithm?
- Line-rate: Highest capacity supported by a communication standard

# Programmability at line-rate



- OpenFlow: Match-Action interface, fixed fields, fixed actions
- P4, RMT, FlexPipe, Xpliant: Protocol-independent match-action pipeline.

# Isn't P4 sufficient?

- Match-action is perfect for forwarding
- But, limiting for stateful algorithms
- Example: RED:

On enqueue:

Calculate average queue size

if  $\min < \text{avg} < \max$

calculate probability  $p$

mark packet with probability  $p$

else if  $\text{avg} > \max$ :

mark packet

# Packet Transactions

- Imperative code block in subset of C (domino) that is atomic and isolated from other such blocks
- One packet transaction per pipeline
- More familiar to NPU, Click programmers

# Programming with Packet Transactions

## Domino

```
#define NUM_FLOWLETS 8000
#define THRESHOLD 5
#define NUM_HOPS 10

struct Packet { int sport; int dport; ...};

int last_time [NUM_FLOWLETS] = {0};
int saved_hop [NUM_FLOWLETS] = {0};

void flowlet(struct Packet pkt) {
    pkt.new_hop = hash3(pkt.sport, pkt.dport, pkt.arrival)
                  % NUM_HOPS;
    pkt.id = hash2(pkt.sport, pkt.dport) % NUM_FLOWLETS;
    if (pkt.arrival - last_time[pkt.id] > THRESHOLD) {
        saved_hop[pkt.id] = pkt.new_hop;
    }
    last_time[pkt.id] = pkt.arrival;
    pkt.next_hop = saved_hop[pkt.id];
}
```

## P4

Stage 1

```
pkt.new_hop = hash3(pkt.sport,
                    pkt.dport,
                    pkt.arrival)
              % NUM_HOPS;
```

```
pkt.id = hash2(pkt.sport, pkt.dport)
          % NUM_FLOWLETS
```

Stage 2

```
pkt.last_time = last_time[pkt.id];
last_time[pkt.id] = pkt.arrival;
```

Stage 3

```
pkt.tmp = pkt.arrival - pkt.last_time;
```

Stage 4


```
pkt.tmp2 = pkt.tmp > 5;
```

Stage 5

```
pkt.saved_hop = saved_hop[pkt.id];
saved_hop[pkt.id] = pkt.tmp2 ?
                    pkt.new_hop :
                    pkt.saved_hop;
```

Stage 6

```
pkt.next_hop = pkt.tmp2 ?
                pkt.new_hop :
                pkt.saved_hop ;
```




# Language constraints on domino

- No loops (for, while, do while)
- No unstructured control flow (goto, break, continue)
- No pointers, heaps

# If Conversion

```
if (pkt.arrival - last_time[pkt.id] > THRESHOLD) {  
    saved_hop [ pkt . id ] = pkt . new_hop ;  
}
```



```
pkt.tmp = pkt.arrival - last_time[pkt.id] > THRESHOLD;  
saved_hop [ pkt . id ] = pkt.tmp  
                        ? pkt . new_hop  
                        : saved_hop [ pkt . id ];
```



# Read and Write Flanks

```
pkt.id = hash2(pkt.sport, pkt.dport) % NUM_FLOWLETS;  
...  
last_time[pkt.id] = pkt.arrival;  
...
```



```
pkt.id = hash2(pkt.sport, pkt.dport) % NUM_FLOWLETS;  
pkt.last_time = last_time[pkt.id]; // Read flank  
...  
pkt.last_time = pkt.arrival;  
...  
last_time[pkt.id] = pkt.last_time; // Write flank
```

# Static Single-Assignment

```
pkt.id = hash2(pkt.sport, pkt.dport) % NUM_FLOWLETS;  
pkt.last_time = last_time[pkt.id];
```

...

```
pkt.last_time = pkt.arrival;  
last_time[pkt.id] = pkt.last_time ;
```



```
pkt.id0 = hash2(pkt.sport, pkt.dport) % NUM_FLOWLETS;  
pkt.last_time0 = last_time[pkt.id0];
```

...

```
pkt.last_time1 = pkt.arrival;
```

...

```
last_time [pkt.id0] = pkt.last_time1 ;
```

# Critical Path Scheduling

```
pkt.id =  
hash2(pkt.sport,  
      pkt.dport)  
% NUM_FLOWLETS
```

```
pkt.last_time = last_time[pkt.id]
```

```
pkt.tmp = pkt.arrival -  
          pkt.last_time
```

```
last_time[pkt.id] = pkt.arrival
```

```
pkt.tmp2 = pkt.tmp >  
           THRESHOLD
```

```
pkt.next_hop =  
hash3(pkt.sport,  
      pkt.dport, pkt.arrival)  
% NUM_HOPS
```

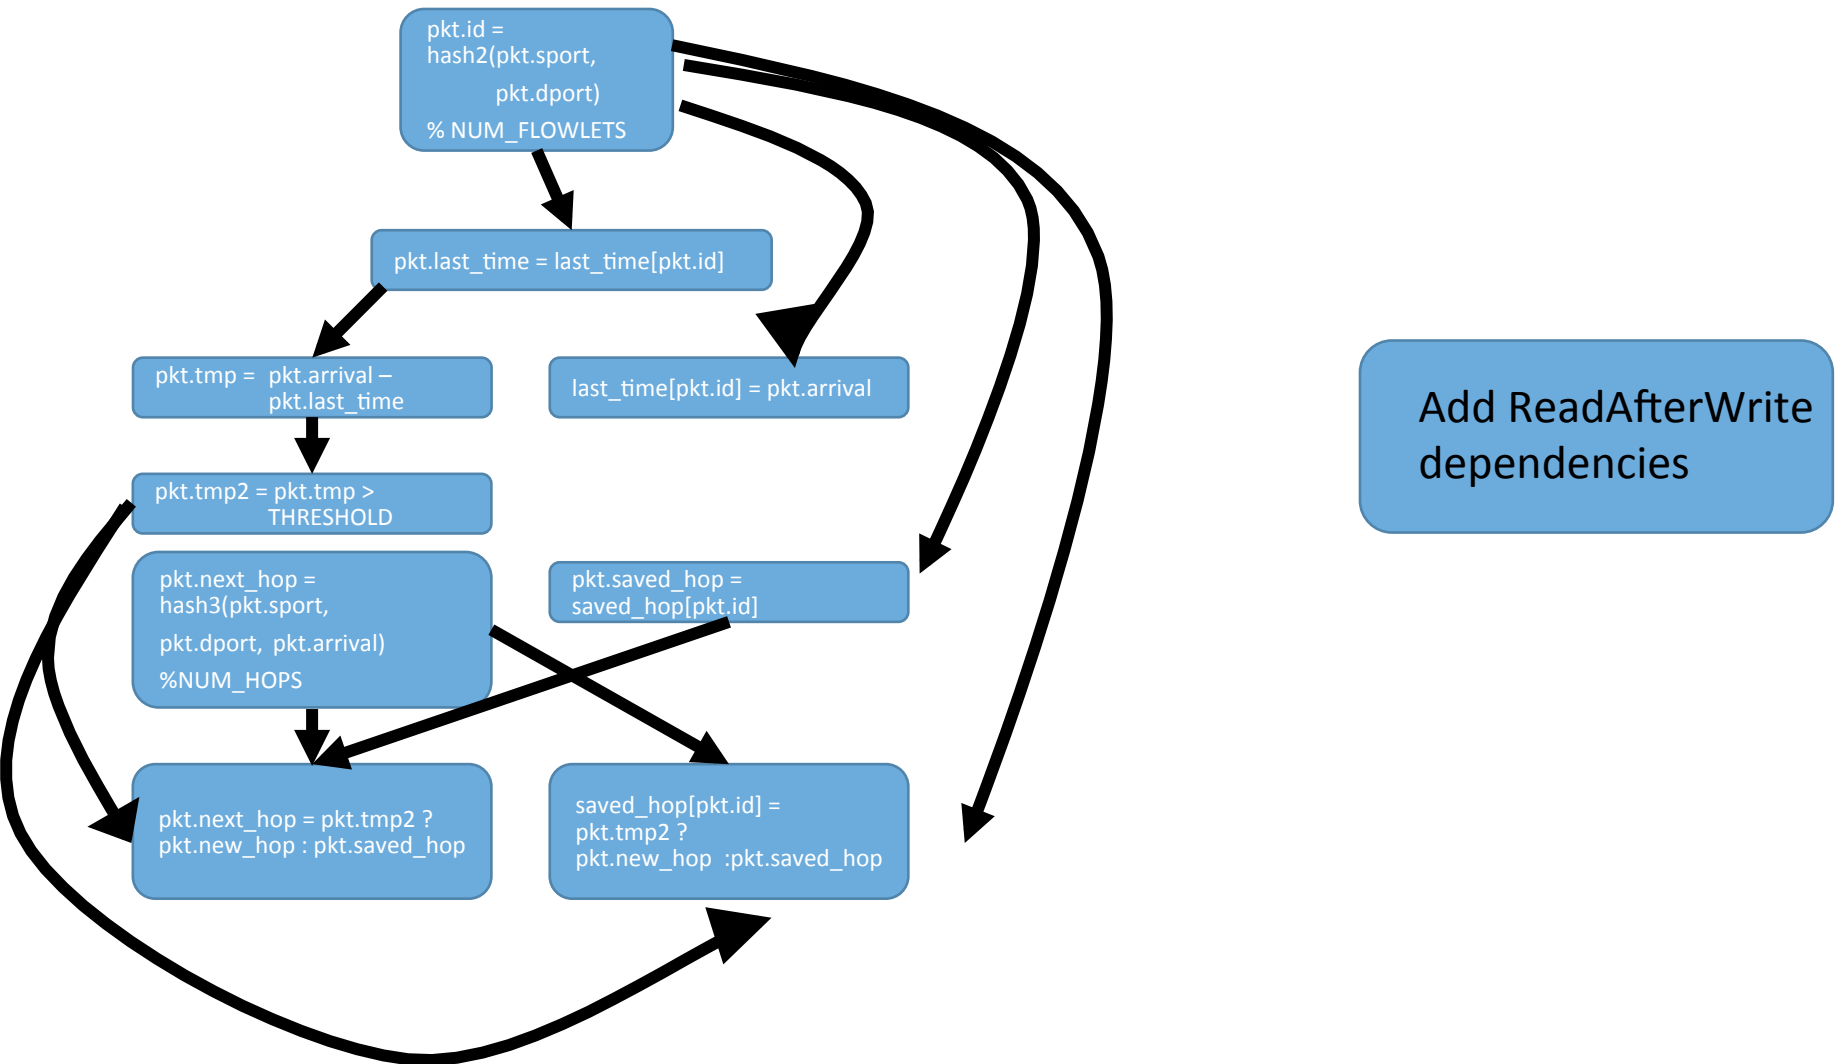
```
pkt.saved_hop =  
saved_hop[pkt.id]
```

```
pkt.next_hop = pkt.tmp2 ?  
pkt.new_hop : pkt.saved_hop
```

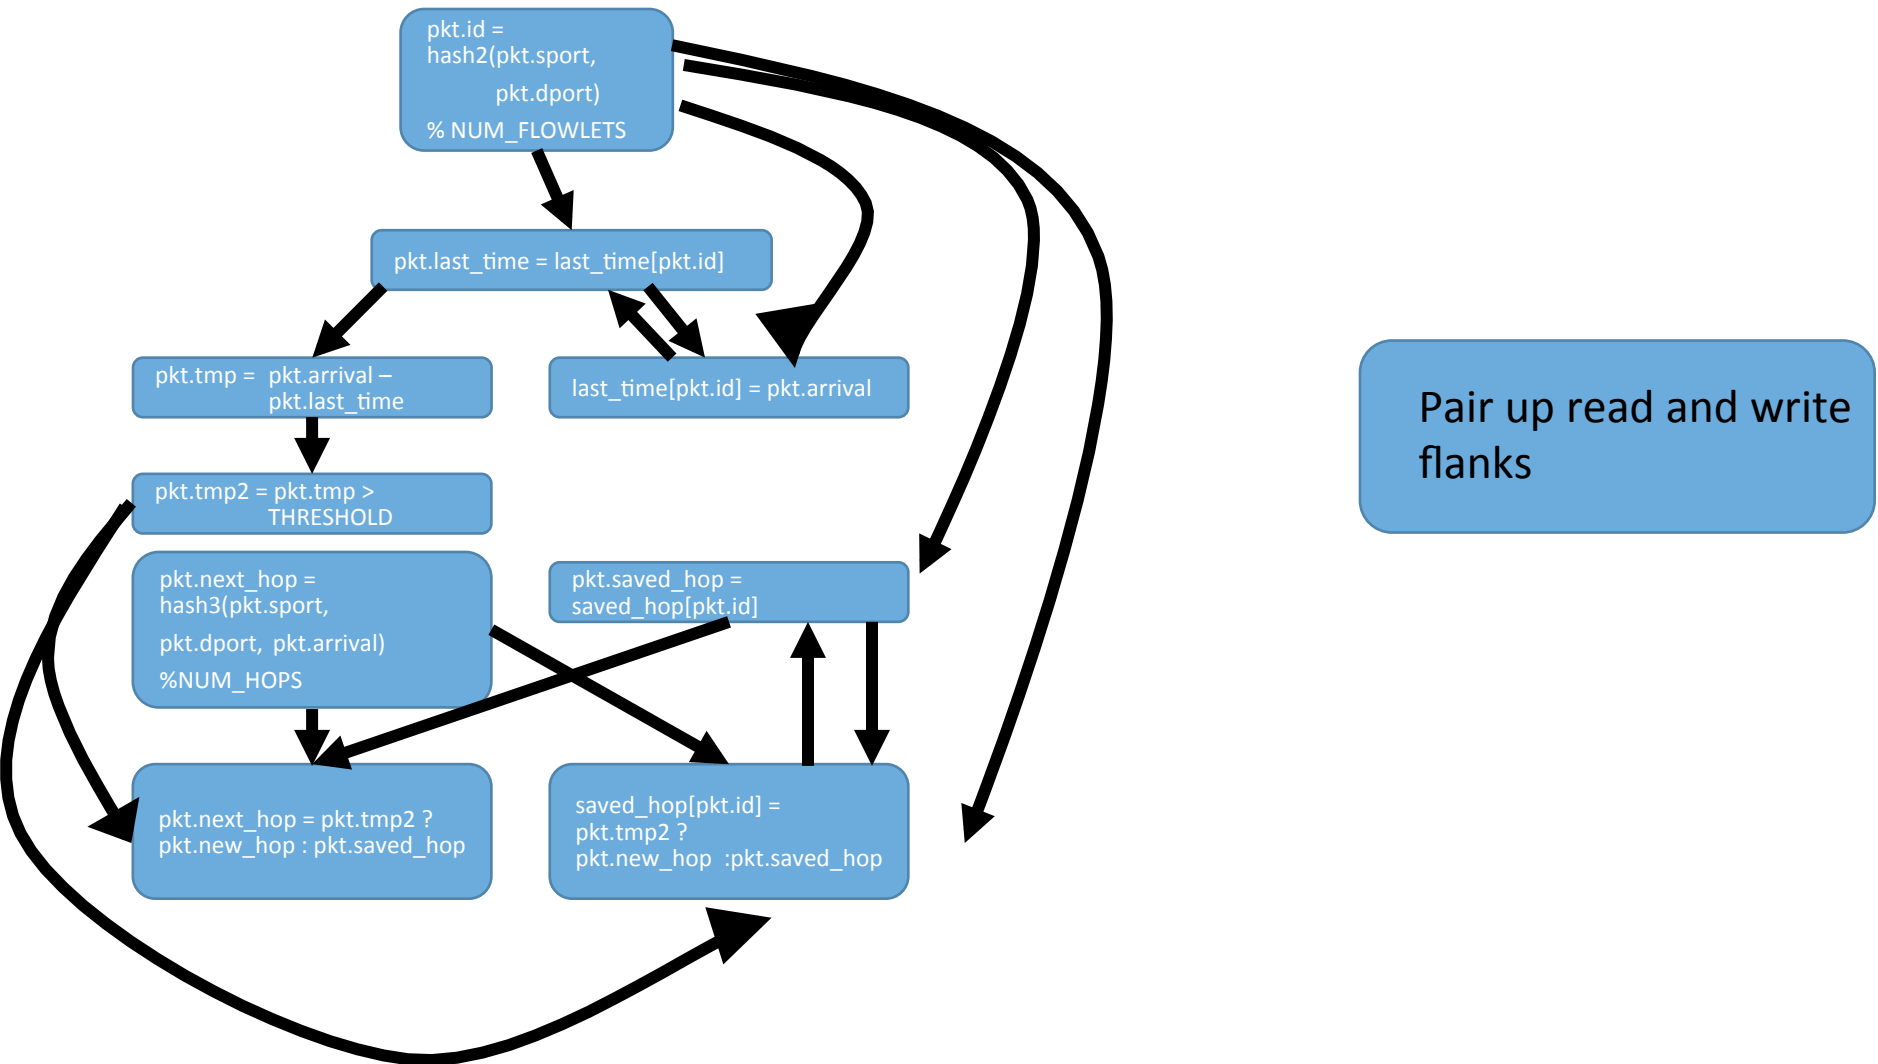
```
saved_hop[pkt.id] =  
pkt.tmp2 ?  
pkt.new_hop : pkt.saved_hop
```

Create one node for  
each instruction.

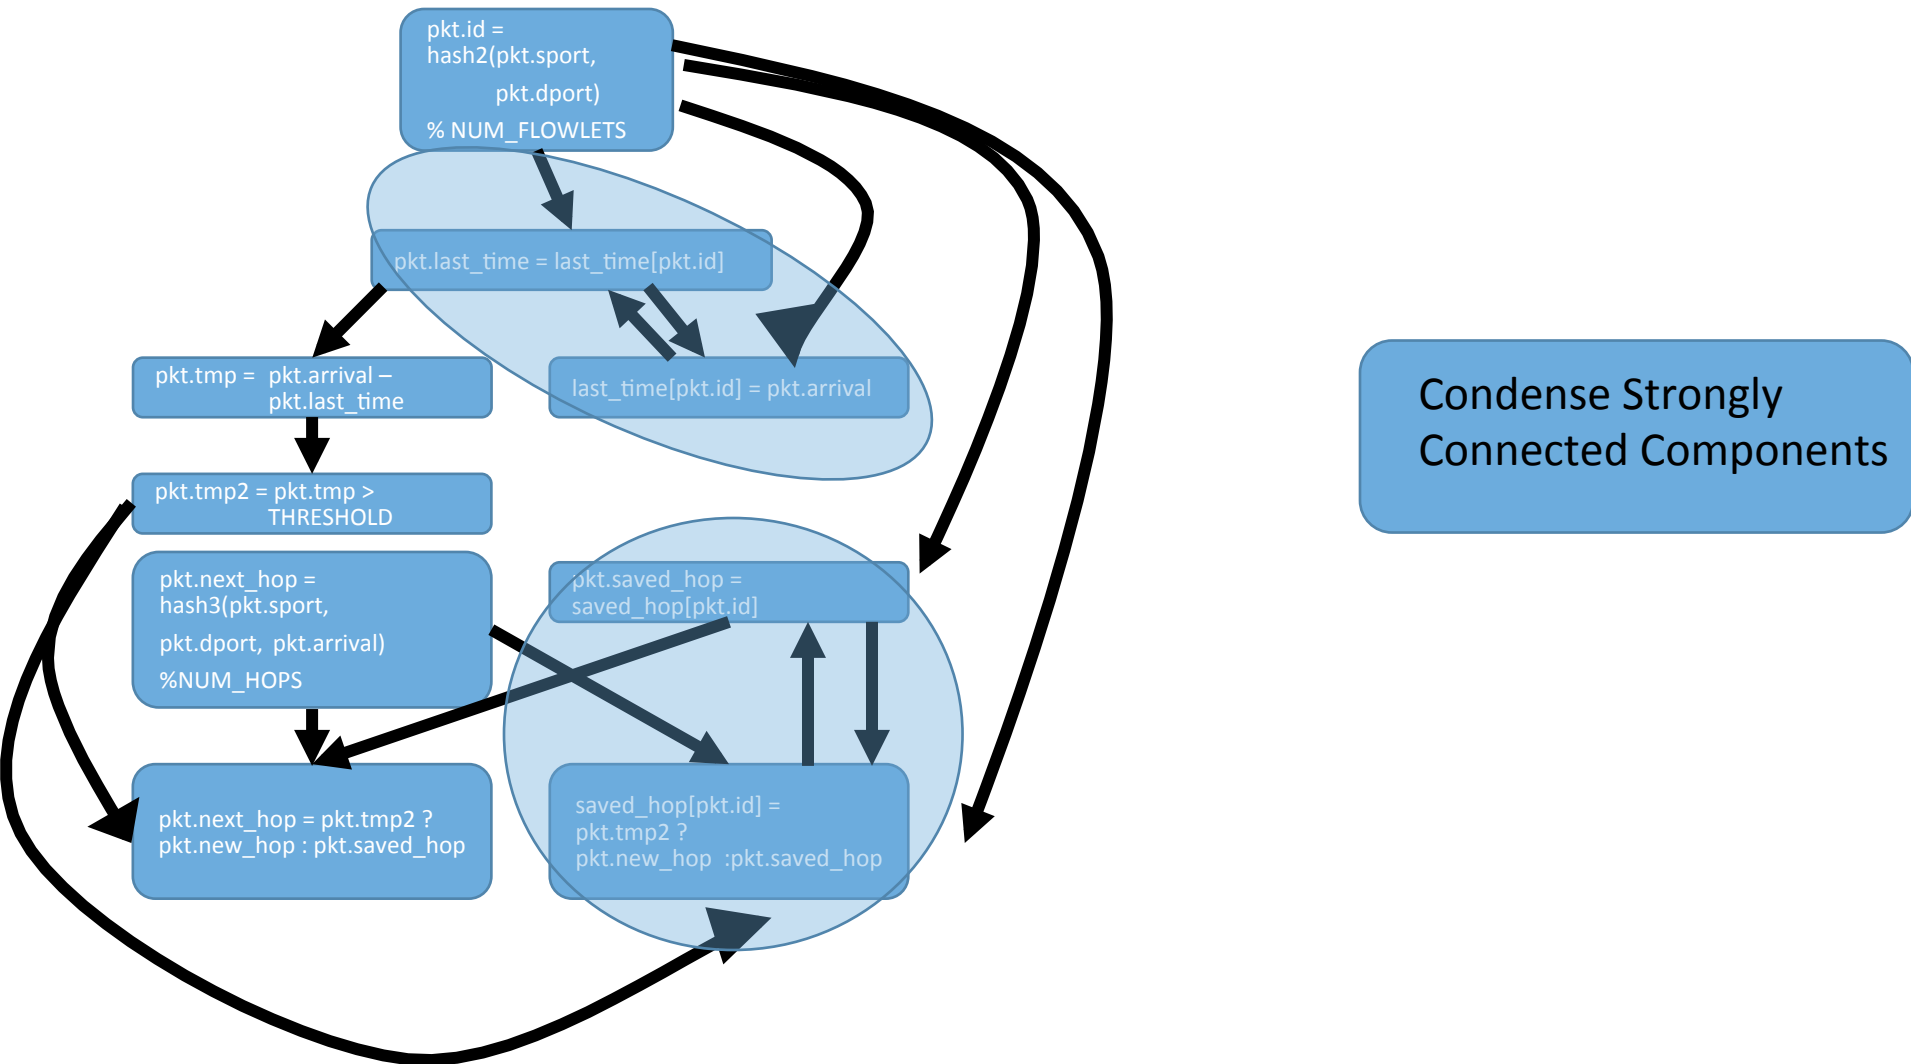
# Critical Path Scheduling



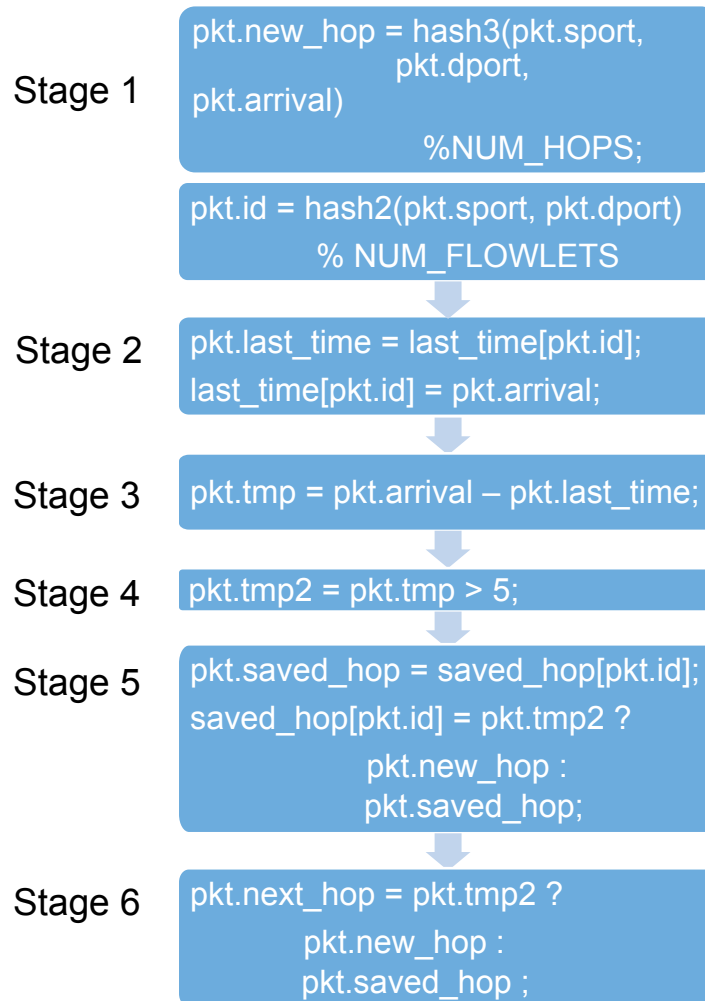
# Critical Path Scheduling



# Critical Path Scheduling



# Critical Path Scheduling



Schedule condensed  
graph

# Generating P4 code

- Required changes to P4
  - Sequential execution semantics
  - Expression support
  - Both available in v1.1
- Encapsulate every SCC in a default action
- Need sequential execution for stateful components
- Thanks to Antonin Bas for help with P4 code generation!



# Checking feasibility at line rate

- So far, haven't constrained action bodies
- Check if action bodies can be mapped to available hardware
- If it can, declare that we can run it at line rate.
  - $x = x + 1$  can be mapped to  $x = x + c$
- If we can't, flag a compiler error E.g.
  - $x = (\text{pkt.y}) ? (x + 1) : x$  can't be mapped to  $x = x + c$ ;

# Closing thoughts

- Constructive proof that we could run a subset of C at line rate
- More familiar abstraction for stateful algorithms (37 LOC for flowlet switching in domino vs 110 in P4)
- Vehicle for higher-level abstractions in packet processing
- Will be open sourcing code soon
- Come check out the demo!