



New Measurement Abstractions for Data Centers

Minlan Yu

University of Southern California

Key Challenges in Data Centers

- Service level agreements
 - Good, predictable performance
- Multi-tenancy
 - Resource isolation across tenants
- Cost reduction
 - Drive the network to high utilization



All require fine-grained network management

Management = *Measurement* + Control

Performance Diagnosis

Measure delay,
throughput of flows

Isolate and fix performance
problems

Security

Capture traffic
anomalies

Filter out or block malicious
traffic/tenants

Traffic Engineering

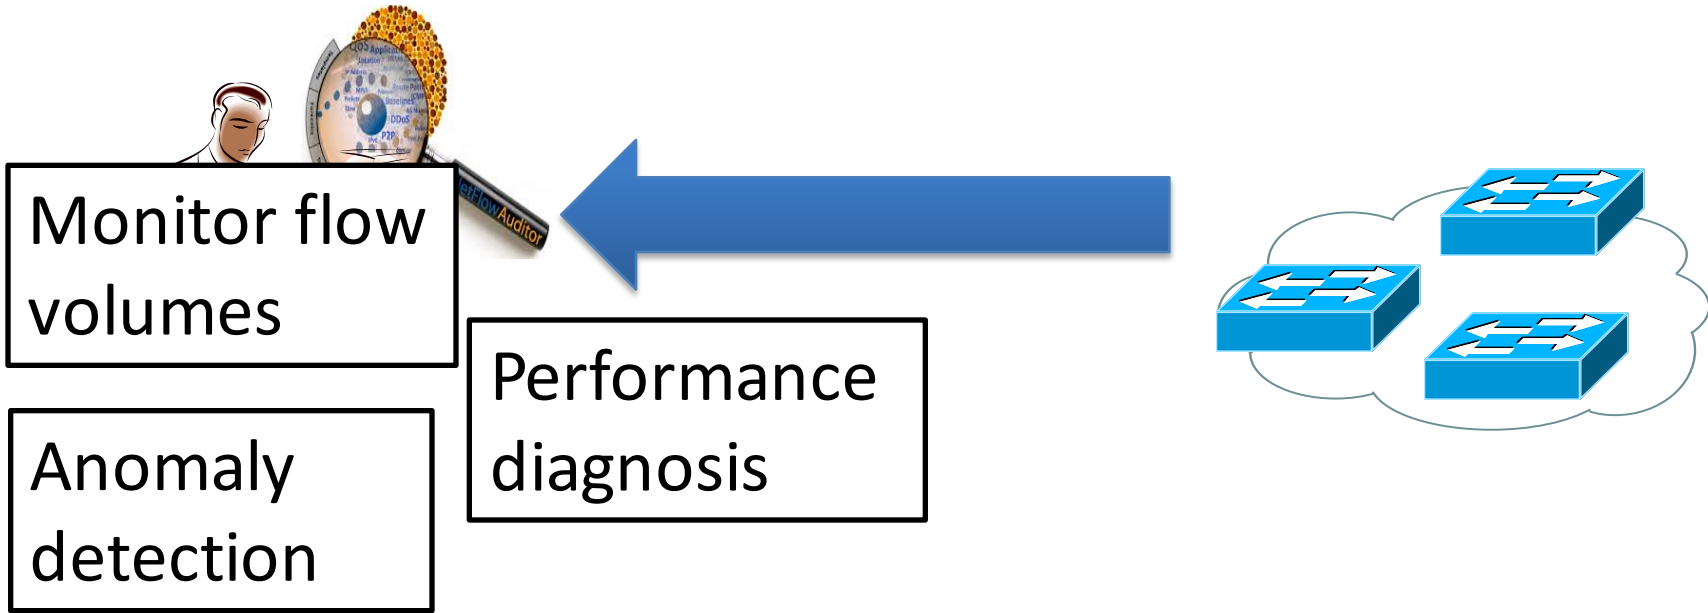
Collect per-flow
statistics and changes

Install routing rules

The state of
today's Datacenter measurement

The need for
new measurement abstractions

Many Measurement Programs



Query independence:

Specify many measurement queries
with simple, programmable abstractions

Many Data Collection Primitives

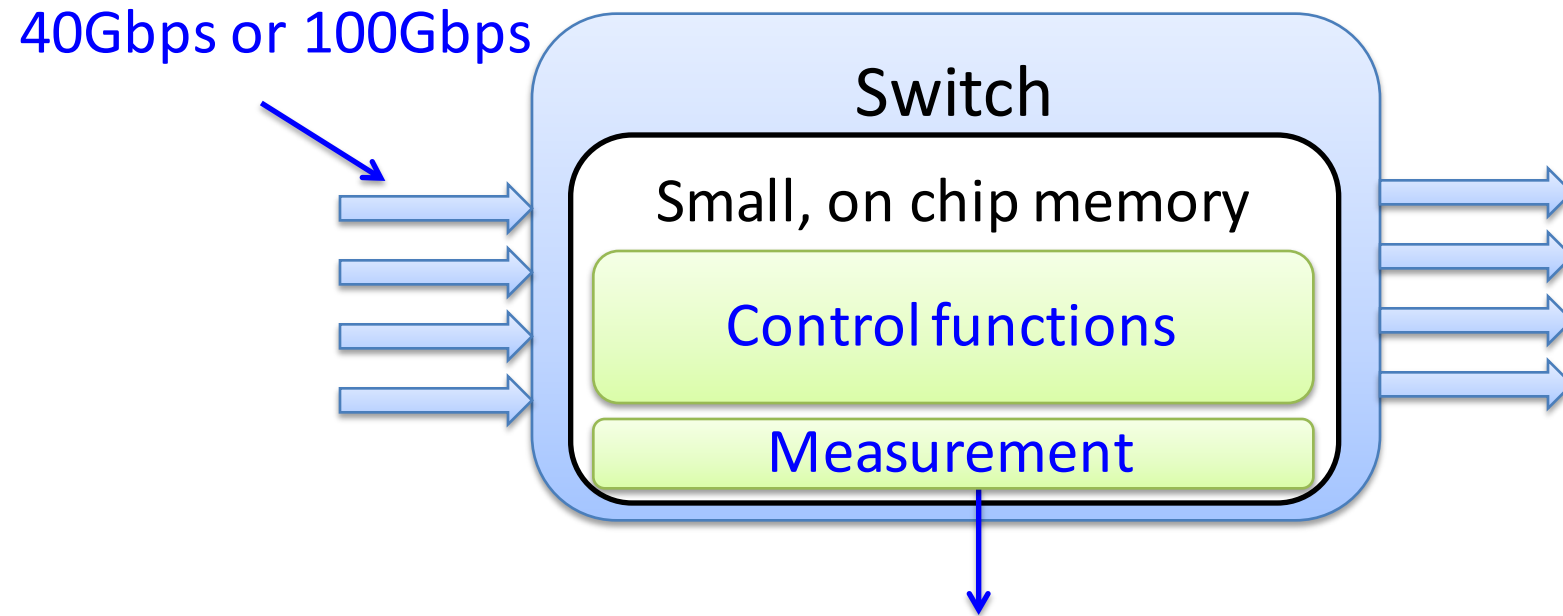


- Passively collect different sources of data
→ Little control on what (not) to measure

Target independence:

Specify which data to collect
independent of where and how it is collected

Limited Measurement Resources



Today, to reduce resource usage

- Low sampling rates
- Miss many important information

Reconfigurable:

Dynamically change the way resources are used for different measurement queries and traffic

A New Abstraction for Measurement



Specify measurement queries


Measurement Framework

Query independence:

Expressive Abstractions for many queries

Reconfigurable: Efficient Runtime


Dynamically configure
devices & manage resources



Automatically collect
the right data

Target independence:

Implementable for many devices



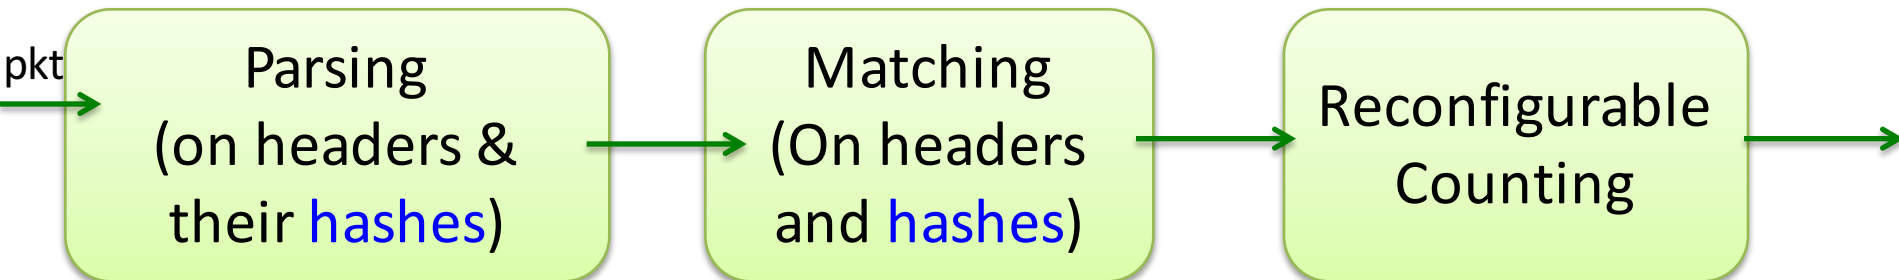
Expressive Abstractions (OpenSketch)

Picking the packets to measure

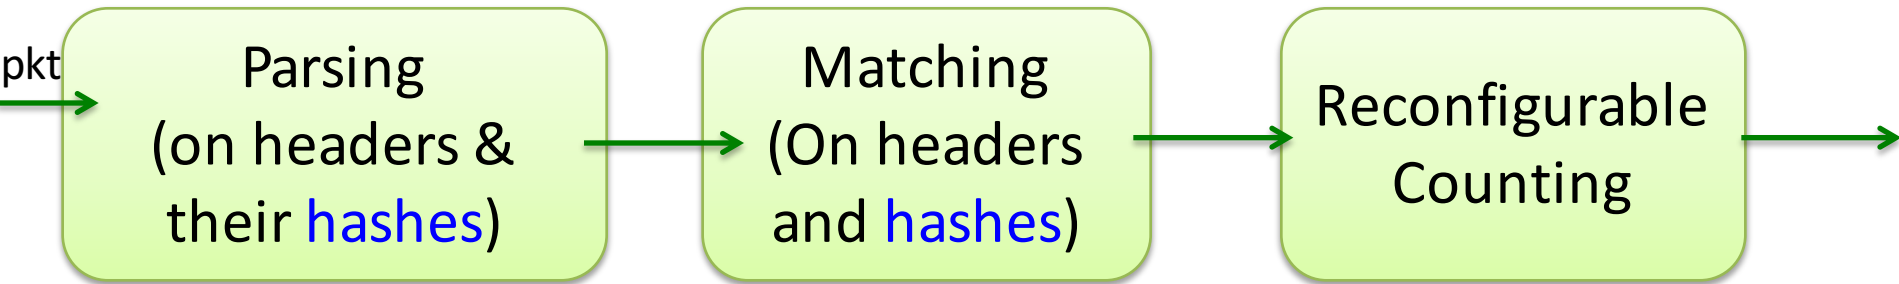
Filter traffic aggregates
(e.g., to host A)
Classify a set of flows
(e.g., a set of malicious source IPs)

Storing & exporting data

Dynamically flows to counters
(e.g., more counters for
elephant flows)




Support Many Queries

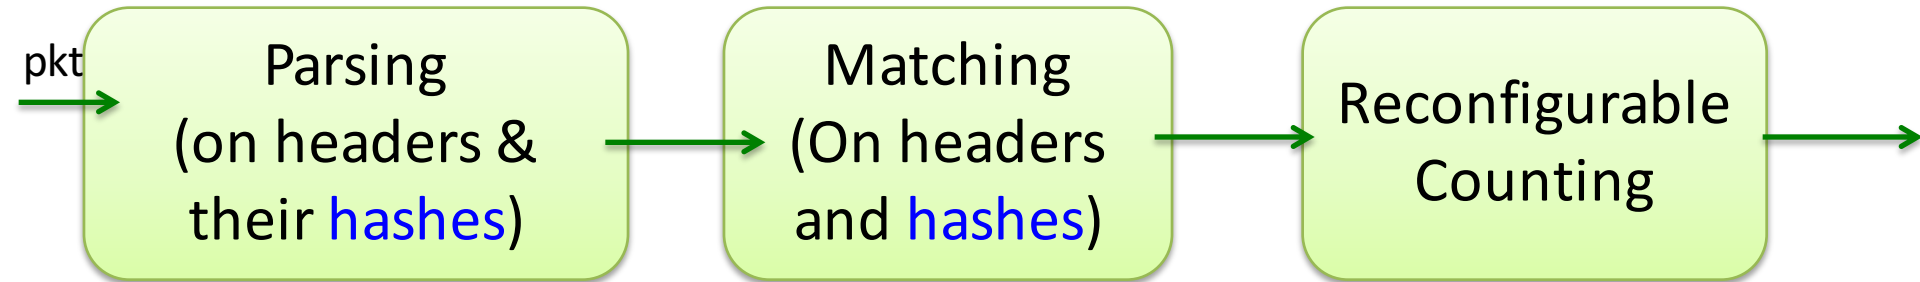


How many unique IPs send traffic to host A?

Bitmap




Support Many Queries



Who's sending a lot to host A?

Count-Min Sketch



Multiple Tables

- Multiple tables for one query
 - Who's sending a lot to host A?
 - *count-min sketch* to count volume of flows
 - *reversible sketch* to identify flows with heavy counts in the count-min sketch



- Multiple tables for many queries
 - Bitmap for counting unique source IPs
 - CountMin sketch for counting traffic for each source IP

Support NetFlow

- NetFlow

- Store flows and their counters in a hash table
- Accounting network utilization, capacity planning, troubleshooting, and attack detection
- It is challenging to implement NetFlow in hardware

SrcIPadd	DstIPadd	Protocol	SrcPort	DstPort	Pkts	Bytes/Pkt
173.100.21.2	10.0.227.12	11	00A2	00A2	11000	1528
173.100.3.2	10.0.227.12	6	15	15	2491	740
173.100.20.2	10.0.227.12	11	00A1	00A1	10000	1428
173.100.6.2	10.0.227.12	6	19	19	2210	1040
173.100.7.2	10.0.227.12	1	41	41	3456	1140
173.100.4.2	10.0.227.12	11	16	16	1929	840
173.100.2.2	10.0.227.12	6	0	0	2228	628
173.100.5.2	10.0.227.12	11	17	17	1325	940

A Better NetFlow

- How to handle hash collisions
 - A larger hash table to reduce collisions
 - Too much memory usage
 - A linked list of buckets, or send to the control plane
 - Too much delay per packet
- Our idea: Embrace hash collisions
 - Encode all the collided (flow, counter) pairs in one bin
 - Network-wide decode in the controller
 - Simple fixed per packet processing
 - Memory usage similar to perfect hashing and actually smaller

A Better Netflow

- Encode Hash a flow into multiple bins

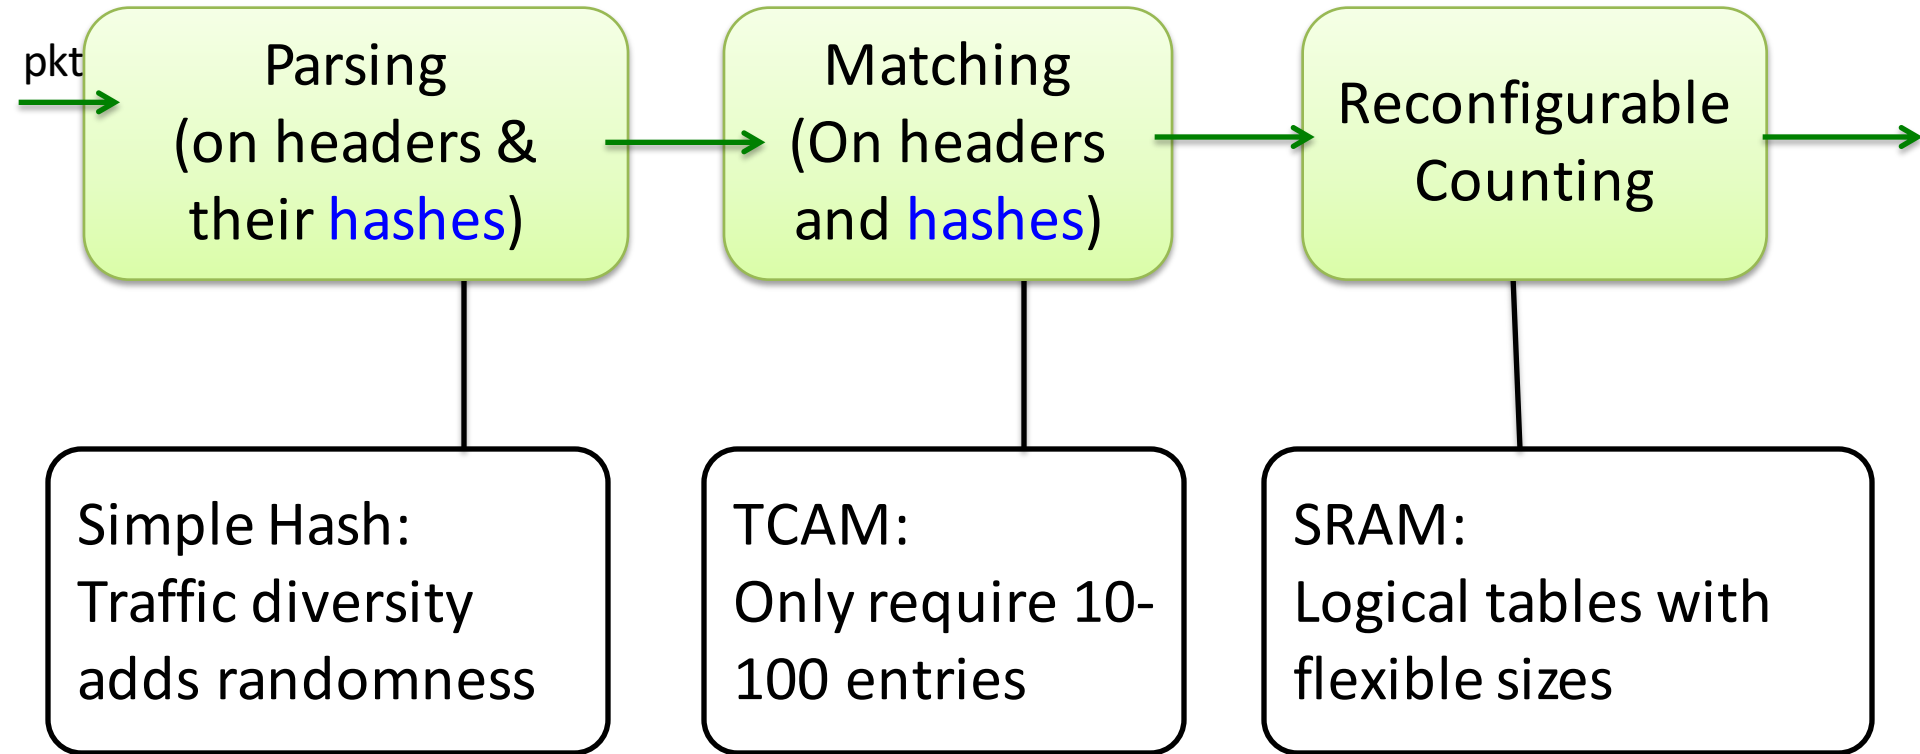
FlowXOR	a	$a \oplus b$	$b \oplus c$...
FlowCounter	1	2	2	
PacketCounter	C_a	C_a+C_b	C_b+C_c	

- Decode
 - Identify bins with a single flow
 - Remove these flows, and then iterate
- Extensions
 - Identify more bins with a single flow
 - Leverage flow information from other hops

Even More Queries...

Measurement Programs	Reconfigurable, hash-based counting table
Loss Detection	Inversible Bloom Filter
Superspreaders	Count-min sketch; Bitmap; Reversible sketch
Traffic change detection	Count-min sketch; Reversible sketch
Traffic entropy on port field	Multi-resolution classifier; Count-min sketch
Flow size distribution	multi-resolution classifier; hash table

Build on P4-compatible Switches



A New Abstraction for Measurement



Specify measurement queries


Measurement Framework

Query independence:

Expressive Abstractions for many queries

Reconfigurable: Efficient Runtime

Dynamically configure
devices & manage resources



Automatically collect
the right data

Target independence:

Implementable for many devices

